

Bibliography

———— Chapter 1 ————

- [Hill '79] Banu Musa brothers (9th century). The book of ingenious devices (Kitab al-hiyal). Translated by D. R. Hill (1979), Springer, p. 44, ISBN 90-277-0833-9.
- [Wiener '48] N. Wiener (1948). Time, communication, and the nervous system. Teleological mechanisms. *Annals of the N.Y. Acad. Sci.* 50 (4): 197-219.
- [Fuegi '03] Fuegi, J., & Francis, J. (2003). Lovelace & Babbage and the creation of the 1843 'notes'. *IEEE Annals of the History of Computing*, 25(4), 16-26.
- [Kagermann '11] Kagermann, H., Lukas, W. D., & Wahlster, W. (2011). Industrie 4.0: Mit dem Internet der Dinge auf dem Weg zur 4. industriellen Revolution. *VDI nachrichten*, 13(1), 2-3.
- [Hey '09] Hey, T., Tansley, S., & Tolle, K. M. (2009). The Fourth Paradigm: Data-Intensive Scientific Discovery (Vol. 1). Redmond, WA: Microsoft Research.
- [Mitchell '97] Mitchell, T.M. (1997) Machine Learning. McGraw Hill.
- [Géron '19] Géron, A. (2019). Hands-on Machine Learning with Scikit-Learn and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems. 2nd Edition, O'Reilly.
- [Bayes '63] Bayes, T. (1763). LII. An Essay Towards Solving a Problem in the Doctrine of Chances. By the late Rev. Mr. Bayes, FRS communicated by Mr. Price, in a letter to John Canton, AMFR S. *Philosophical transactions of the Royal Society of London*, (53), 370-418.
- [Provost '13] Provost, F., & Fawcett, T. (2013). Data Science for Business: What you need to know about data mining and data-analytic thinking. " O'Reilly Media, Inc."
- [Russell '21] Russell, S., Norvig, P. (2021). Artificial intelligence: A Modern Approach (4th Edition). Prentice Hall.
- [Chapmann '17] Chapmann, J. (2017) Machine Learning: Fundamental Algorithms for Supervised and Unsupervised Learning With Real-World Applications.
- [Knuth '97] Knuth, D. E. (1997). The Art of Computer Programming (Vol. 1). Pearson Education.
- [Bhargava '16] Bhargava, A. (2016). Grokking Algorithms: An illustrated guide for programmers and other curious people. Manning Publications Co.
- [Han '11] Han, J., Pei, J., & Kamber, M. (2011). *Data Mining: Concepts and Techniques*. 3rd Edition. Elsevier.
- [Allen '83] Allen, J. F. (1983). Maintaining knowledge about temporal intervals. *Communications of the ACM*, 26(11), 832-843.
- [Schneider '19] Schneider, S., Baevski, A., Collobert, R., & Auli, M. (2019). wav2vec: Unsupervised pre-training for speech recognition. arXiv preprint arXiv:1904.05862.

———— Chapter 2 ————

- [Wolpert '97] Wolpert, D.H., Macready, W.G. (1997), "No Free Lunch Theorems for Optimization", *IEEE Transactions on Evolutionary Computation* 1, 67
- [Topol '19] Topol, E. (2019). Deep medicine: how artificial intelligence can make healthcare human again. Hachette UK.
- [Zheng '18] Zheng, A., & Casari, A. (2018). Feature Engineering for Machine Learning: Principles and Techniques for Data Scientists. " O'Reilly Media, Inc."
- [Deisenroth '19] Deisenroth, M.P, Faisal, A. and Ong C.S. (2019) Mathematics for Machine Learning. Cambridge University Press.
- [Saul '00] Saul, L. K., & Roweis, S. T. (2000). An introduction to locally linear embedding. unpublished. Available at: <http://www.cs.toronto.edu/~roweis/lle/publications.html>.
- [Van der Maaten '08] Van der Maaten, L.J.P.; Hinton, G.E. (2008). Visualizing Data Using t-SNE. *Journal of Machine Learning Research*. 9: 2579–2605.

- [McInnes '18] McInnes, L., Healy, J., & Melville, J. (2018). Umap: Uniform manifold approximation and projection for dimension reduction. *arXiv preprint arXiv:1802.03426*.
- [Barnes '86] Barnes, J. & Hut, P. (1986). A hierarchical $O(N \log N)$ force-calculation algorithm. *Nature*. 324 (4): 446–449
- [Harrison '19] Harrison M. (2019) Machine Learning Pocket Reference: Working with Structured Data in Python. O'Reilly Media.

———— Chapter 3 ————

- [Griffiths '09] Griffin Dawn (2009) *Head First Statistics*. Publisher: O'Reilly.
- [Javaheri '21] Javaheri, T., Homayounfar, M., Amoozgar, Z., Reiazi, R., Homayounieh, F., Abbas, E., ... & Rawassizadeh, R. (2021). CovidCTNet: an open-source deep learning approach to diagnose covid-19 using small cohort of CT images. *NPJ digital medicine*, 4(1), 1-10.
- [Rawassizadeh '19] Rawassizadeh, R., Dobbins, C., Akbari, M. & Pazzani, M. (2019) Indexing Multivariate Data through Spatio -Temporal Event Detection and Clustering. *Sensors*, 19 (3), 448.
- [Rawassizadeh '13] Rawassizadeh, R., Tomitsch, M., Wac, K., & Tjoa, A. M. (2013). UbiqLog: a generic mobile phone-based life-log framework. *Personal and ubiquitous computing*, 17(4), 621-637.
- [Tufte '01] Tufte, E. R. (2001). *The Visual Display of Quantitative Information* (Vol. 2). Cheshire, CT: Graphics press.
- [Abela '13] Abela, A. (2013) *Advanced Presentations by Design: Creating Communication that Drives Action*. Pfeiffer, 2nd edition.
- [Knafllic '15] Knafllic, C. N. (2015). *Storytelling with data: A data visualization guide for business professionals*. John Wiley & Sons.
- [Wilke '19] Wilke, C. O. (2019). *Fundamentals of Data Visualization: A Primer on Making Informative and Compelling Figures*. O'Reilly Media.
- [Rawassizadeh '15] Rawassizadeh, R., Momeni, E., Dobbins, C., Mirza-Babaei, P., & Rahnamoun, R. (2015). Lesson learned from collecting quantified self information via mobile and wearable devices. *Journal of Sensor and Actuator Networks*, 4(4), 315-335.
- [Rumsey '15] Rumsey, D. J. (2015). *U Can: statistics for dummies*. John Wiley & Sons.
- [Rumsey '09] Rumsey, D. J. (2009). *Statistics II for Dummies*. Wiley Publishing, Inc.
- [Boslaugh '12] Boslaugh, S. (2012). *Statistics in a nutshell: A desktop quick reference*. " O'Reilly Media, Inc."
- [Skiena '17] Skiena, S. S. (2017). *The Data Science Design Manual*. Springer.
- [Neyman '37] Neyman, J. (1937). Outline of a theory of statistical estimation based on the classical theory of probability. *Philosophical Transactions of the Royal Society of London. Series A, Mathematical and Physical Sciences*, 236(767), 333-380.
- [Chandola '09] Chandola, V., Banerjee, A., & Kumar, V. (2009). Anomaly detection: A survey. *ACM computing surveys (CSUR)*, 41(3), 15.
- [Aggarwal '01] Aggarwal, C. C., Hinneburg, A., & Keim, D. A. (2001). On the surprising behavior of distance metrics in high dimensional space. In *International conference on database theory* (pp. 420-434). Springer, Berlin, Heidelberg.
- [Goodfellow '16] Goodfellow, I., Bengio, Y., Courville, A., & Bengio, Y. (2016). *Deep learning* (Vol. 1). Cambridge: MIT press.
- [Shapiro '65] Shapiro, S. S.; Wilk, M. B. (1965). "An analysis of variance test for normality (complete samples)". *Biometrika*. 52 (3–4): 591–611.
- [Bonferroni '36] Bonferroni, C. E., Teoria statistica delle classi e calcolo delle probabilità, *Pubblicazioni del R Istituto Superiore di Scienze Economiche e Commerciali di Firenze* 1936
- [Tukey '49] Tukey, J. W. (1949). Comparing individual means in the analysis of variance. *Biometrics*, 99-114.
- [Sullivan '12] Sullivan, G. M., & Feinn, R. (2012). Using Effect Size—or Why the P-value is not Enough. *Journal of graduate medical education*, 4(3), 279-282.
- [Cohen '92] Cohen, J. (1992). A power primer. *Psychological bulletin*, 112(1), 155 - 159.

- [Cliff '93] Cliff, Norman (1993). "Dominance Statistics: Ordinal Analyses to Answer Ordinal Questions". *Psychological Bulletin*. 114 (3): 494-509.
- [Stigler '92] Stigler, S. M. (1989). "Francis Galton's Account of the Invention of Correlation". *Statistical Science*. 4 (2), 73-79.
- [Spearman '04] Spearman, C. (1904). The proof and measurement of association between two things. *The American journal of psychology*, 15(1), 72-101.
- [Kendall '38] Kendall, M. G. (1938). A new measure of rank correlation. *Biometrika*, 30(1/2), 81-93.
- [Shannon '49] Shannon, C. E. (1949). Communication theory of secrecy systems. *Bell system technical journal*, 28(4), 656-715.
- [Bishop '06] Bishop, C.M. (2006) Pattern Recognition and Machine Learning. Springer.
- [Kullback '51] Kullback, S., Leibler, R.A. (1951). On information and sufficiency. *Annals of Mathematical Statistics*. 22 (1): 79-86.
- [Bruce '17] Bruce, P., & Bruce, A. (2017). Practical Statistics for Data Scientists: 50 Essential Concepts. O'Reilly Media, Inc.

———— Chapter 4 ————

- [Dunn '61] Dunn, O.J. (1961). Multiple Comparisons Among Means. *Journal of the American Statistical Association*. 56 (293), 52-64
- [Aggarwal '15] Aggarwal, C. C. (2015). *Data Mining: The Textbook*. Springer.
- [Rao '00] Rao, K. R., & Yip, P. C. (2000). *The Transform and Data Compression Handbook*. CRC press.
- [Pighizzini '01] Pighizzini, G. (2001). How hard is computing the edit distance?. *Information and Computation*, 165(1), 1-13.
- [Ratanamahatana '05] Ratanamahatana, C. A., & Keogh, E. (2005). Three Myths about Dynamic Time Warping Data Mining. In *Proceedings of the 2005 SIAM International Conference on Data Mining* (pp. 506-510). Society for Industrial and Applied Mathematics.
- [Xu '05] Xu, R., & Wunsch, D. (2005). Survey of Clustering Algorithms. *IEEE Transactions on neural networks*, 16(3), 645-678.
- [Ester '96] Ester, M., Kriegel, H. P., Sander, J., & Xu, X. (1996). A density-based algorithm for discovering clusters in large spatial databases with noise. In *KDD* (Vol. 96, No. 34, pp. 226-231).
- [Ankerst '99] Ankerst, M., Breunig, M. M., Kriegel, H. P., & Sander, J. (1999). OPTICS: ordering points to identify the clustering structure. In *ACM Sigmod record* (Vol. 28, No. 2, pp. 49-60). ACM.
- [Sibson '73] Sibson, R. (1973). SLINK: An Optimally Efficient Algorithm for the Single-Link Cluster Method. *The Computer Journal*. British Computer Society. 16 (1): 30-34.
- [Defays '77] Defays, D. (1977). An Efficient Algorithm for a Complete Link Method. *The Computer Journal*. British Computer Society. 20 (4): 364-366.
- [Kaufman '09] Kaufman, L., & Rousseeuw, P. J. (2009). Finding groups in data: an introduction to cluster analysis (Vol. 344). John Wiley & Sons.
- [Larose '15] Larose, D. T. (2015). Data Mining and Predictive Analytics. John Wiley & Sons.
- [Manning '10] Manning, C., Raghavan, P., & Schütze, H. (2010). Introduction to information retrieval. *Natural Language Engineering*, 16(1), 100-103.
- [Kurzweil '06] Kurzweil, R. (2006). The Singularity is Near: When Humans Transcend Biology. Gerald Duckworth & Co.
- [Deza '09] Deza, M. M., & Deza, E. (2009). Encyclopedia of distances. In *Encyclopedia of distances* (pp. 1-583). Springer, Berlin, Heidelberg.
- [Zhang '96] Zhang, T., Ramakrishnan, R., & Livny, M. (1996). BIRCH: an efficient data clustering method for very large databases. In *ACM Sigmod Record* (Vol. 25, No. 2, pp. 103-114).
- [Guha '98] Guha, S., Rastogi, R., & Shim, K. (1998). CURE: An Efficient Clustering Algorithm for Large Databases. In *ACM Sigmod Record* (Vol. 27, No. 2, pp. 73-84).
- [Edgeworth '1908] Edgeworth, F. Y. (1908). "On the / errors of frequency-constants". *Journal of the Royal Statistical Society*. 71 (3): 499-512. doi:10.2307/2339293.

- [Duda '73] Duda, R. O., & Hart, P. E. (1973). Pattern classification and scene analysis. *A Wiley-Interscience Publication, New York: Wiley, 1973*.
- [Dempster '77] Dempster, A.P.; Laird, N.M.; Rubin, D.B. (1977). "Maximum Likelihood from Incomplete Data via the EM Algorithm". *Journal of the Royal Statistical Society, Series B*. 39 (1): 1–38.
- [Yule '36] Yule, G. U., Filon, L. N. G. (1936). "Karl Pearson. 1857-1936". *Obituary Notices of Fellows of the Royal Society*. 2 (5): 72
- [Zadeh '96] Zadeh, L. A. (1996). Fuzzy sets. In *Fuzzy sets, fuzzy logic, and fuzzy systems: selected papers by Lotfi A Zadeh* (pp. 394-432).
- [Rawassizadeh '19] Rawassizadeh, R., Dobbins, C., Akbari, M., & Pazzani, M. (2019). Indexing multivariate mobile data through spatio-temporal event detection and clustering. *Sensors*, 19(3), 448.
- [McInnes '17] McInnes, L., Healy, J., & Astels, S. (2017). hdbscan: Hierarchical density based clustering. *J. Open Source Softw.*, 2(11), 205.
- [Leskovec '14] Leskovec, J., Rajaraman, A., & Ullman, J. D. (2014). Mining of Massive Datasets. *Cambridge University Press*.
- [Rousseeuw '87] Rousseeuw, P. J. (1987). Silhouettes: a graphical aid to the interpretation and validation of cluster analysis. *Journal of computational and applied mathematics*, 20, 53-65.
- [Dunn '74] Dunn, J. C. (1974). Well-separated clusters and optimal fuzzy partitions. *Journal of cybernetics*, 4(1), 95-104.
- [Davies '79] Davies, D. L., & Bouldin, D. W. (1979). A cluster separation measure. *IEEE transactions on pattern analysis and machine intelligence*, (2), 224-227.
- [Caliński '74] Caliński, T., & Harabasz, J. (1974). A dendrite method for cluster analysis. *Communications in Statistics-theory and Methods*, 3(1), 1-27.

———— Chapter 5 ————

- [Zaki '14] Zaki, M. J., Meira Jr, W., & Meira J., W. (2014). Data Mining and Analysis: Fundamental Concepts and Algorithms. Cambridge University Press.
- [Sedgewick '13] Sedgewick, R., & Flajolet, P. (2013). An introduction to the analysis of algorithms. Pearson Education India.
- [McDowell '16] McDowell, G. L. (2016). Cracking the Coding Interview: 189 Programming Questions and Solutions. CareerCup, LLC.
- [Broder '97] Broder, A. Z. (1997). On the resemblance and containment of documents. In *Proceedings. Compression and Complexity of SEQUENCES 1997* (Cat. No. 97TB100171) (pp. 21-29).
- [Comer '79] Comer, D. (1979). Ubiquitous B-tree. *ACM Computing Surveys (CSUR)*, 11(2), 121-137.
- [Bayer '70] Bayer, R., & McCreight, E. (1970). Organization and maintenance of large ordered indices. In *Proceedings of the 1970 ACM SIGFIDET (Now SIGMOD) Workshop on Data Description, Access and Control* (pp. 107-141).
- [Coulom '07] Coulom, R. (2007). Efficient selectivity and backup operators in Monte-Carlo tree search. In *Computers and Games: 5th International Conference, CG 2006, Turin, Italy, May 29-31, 2006. Revised Papers 5* (pp. 72-83). Springer Berlin Heidelberg.
- [Silver '18] Silver, D., Hubert, T., Schrittwieser, J., Antonoglou, I., Lai, M., Guez, A., ... & Hassabis, D. (2018). A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play. *Science*, 362(6419), 1140-1144.
- [Bloom '70] Bloom, B. H. (1970). Space/time trade-offs in hash coding with allowable errors. *Communications of the ACM*, 13(7), 422-426.
- [Pugh '90] Pugh, W. (1990). Skip lists: a probabilistic alternative to balanced trees. *Communications of the ACM*, 33(6), 668-676.
- [Agrawal '94] Agrawal, R., & Srikant, R. (1994). Fast Algorithms for Mining Association Rules. In *Proceedings of 20th International Conference of Very Large Data Bases, VLDB* (Vol. 1215, pp. 487-499).
- [Park '95] Park, J. S., Chen, M. S., & Yu, P. S. (1995). An effective hash-based algorithm for mining association rules. *Acm sigmod record*, 24(2), 175-186.

- [Han '00] Han, J., Pei, J., & Yin, Y. (2000). Mining Frequent Patterns Without Candidate Generation. In *ACM Sigmod Record* (Vol. 29, No. 2, pp. 1-12).
- [Zaki '97] Zaki, M. J., Parthasarathy, S., Ogihara, M., & Li, W. (1997). New Algorithms for Fast Discovery of Association Rules. In *Knowledge Discovery and Data Mining* (Vol. 97, pp. 283-286).
- [Srikant '96] Srikant, R., & Agrawal, R. (1996). Mining sequential patterns: Generalizations and performance improvements. In *International Conference on Extending Database Technology* (pp. 1-17).
- [Zaki '01] Zaki, M. J. (2001). SPADE: An Efficient Algorithm for Mining Frequent Sequences. *Machine Learning*, 42(1-2), 31-60.
- [Han '00] Han, J., Pei, J., Mortazavi-Asl, B., Chen, Q., Dayal, U., & Hsu, M. C. (2000). FreeSpan: Frequent Pattern-Projected Sequential Pattern Mining. In *Proceedings of the sixth ACM SIGKDD international conference on Knowledge discovery and data mining* (pp. 355-359).
- [Han '01] Han, J., Pei, J., Mortazavi-Asl, B., Pinto, H., Chen, Q., Dayal, U., & Hsu, M. (2001, April). Prefixspan: Mining sequential patterns efficiently by prefix-projected pattern growth. In *Proceedings of the 17th International Conference on Data Engineering* (pp. 215-224).
- [Han '05] Han, J., Pei, J., & Yan, X. (2005). Sequential pattern mining by pattern-growth: Principles and extensions. In *Foundations and Advances in Data Mining* (pp. 183-220). Springer, Berlin, Heidelberg.
- [Aggarwal '14] Aggarwal, C. C., & Han, J. (Eds.). (2014). *Frequent Pattern Mining*. Springer.
- [Pei '04] Pei, J., Han, J., Mortazavi-Asl, B., Pinto, H., Chen, Q., Dayal, U., & Hsu, M. C. (2004). Mining sequential patterns by pattern-growth: The prefixspan approach. *IEEE Transactions on Knowledge & Data Engineering*, (11), 1424-1440.
- [Gonzalez '17] Gonzalez, R. C., Woods, R. E., & Eddins, S. L. (2017). *Digital Image Processing, 4th Edition*. Pearson-Prentice-Hall.
- [Jurafsky '18] Jurafsky, D. and H. Martin, J. (2018) Speech and Language Processing (3rd edition draft) online: <https://web.stanford.edu/~jurafsky/slp3/>
- [Baum '70] Baum, L. E., Petrie, T., Soules, G., & Weiss, N. (1970). A maximization technique occurring in the statistical analysis of probabilistic functions of Markov chains. *The annals of mathematical statistics*, 41(1), 164-171.
- [Baum '67] Baum, L. E., & Eagon, J. A. (1967). An inequality with applications to statistical estimation for probabilistic functions of Markov processes and to a model for ecology. *Bulletin of the American Mathematical Society*. 360-363.

———— Chapter 6 ————

- [Ozdemir '18] Ozdemir, S., & Susarla, D. (2018). Feature Engineering Made Easy: Identify unique features from your dataset in order to build powerful machine learning systems. Packt Publishing Ltd.
- [Chandrashekar '14] Chandrashekar, G., & Sahin, F. (2014). A Survey on Feature Selection Methods. *Computers & Electrical Engineering*, 40(1), 16-28.
- [Quinlan '86] Quinlan, J. R. (1986). Induction of decision trees. *Machine learning*, 1(1), 81-106.
- [Tibshirani '96] Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society: Series B (Methodological)*, 58(1), 267-288.
- [Rawassizadeh '15] Rawassizadeh, R., Price, B. A., & Petre, M. (2015). Wearables: Has the Age of Smartwatches Finally Arrived?. *Communications of the ACM*, 58(1), 45-47.
- [Menon '20] Menon, S., Damian, A., Hu, S., Ravi, N., & Rudin, C. (2020). Pulse: Self-supervised photo upsampling via latent space exploration of generative models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition* (pp. 2437-2445).
- [Moody '88] Moody, J. (1988). Fast learning in multi-resolution hierarchies. *Advances in neural information processing systems*.
- [Box '64] Box, G. E. P. and Cox, D. R. (1964). An Analysis of Transformations, *Journal of the Royal Statistical Society, Series B*, 26, 211-252.
- [Ng '18] Ng, A. (2018). Machine Learning Yearning. URL: [http://www.mlyearning.org/\(96\)](http://www.mlyearning.org/(96)).
- [Mikolov '13] Mikolov, T., Chen, K., Corrado, G., & Dean, J. (2013). Efficient Estimation of Word Representations in Vector Space. *arXiv preprint arXiv:1301.3781*.

- [Pennington '14] Pennington, Jeffrey, Socher, Richard, and Manning, Christopher D. GloVe: Global vectors for word representation. In Proceedings of the 2014 conference on empirical methods in natural language processing, pp. 1532–1543, 2014.
- [Bojanowski '17] Bojanowski, P., Grave, E., Joulin, A., & Mikolov, T. (2017). Enriching word vectors with subword information. *Transactions of the Association for Computational Linguistics*, 5, 135-146.
- [Rose '10] Rose, S., Engel, D., Cramer, N., & Cowley, W. (2010). Automatic keyword extraction from individual documents. *Text mining: applications and theory*, 1-20.
- [Mihalcea '04] Mihalcea, R., & Tarau, P. (2004). TextRank: Bringing order into text. In Proceedings of the 2004 conference on empirical methods in natural language processing (pp. 404-411).
- [Campos '18] Campos, R., Mangaravite, V., Pasquali, A., Jorge, A. M., Nunes, C., & Jatowt, A. (2018). Yake! collection-independent automatic keyword extractor. In *Advances in Information Retrieval: 40th European Conference on IR Research, ECIR 2018, Grenoble, France, March 26-29, 2018, Proceedings 40* (pp. 806-810). Springer International Publishing.
- [Ji '12] Ji, S., Xu, W., Yang, M., & Yu, K. (2012). 3D convolutional neural networks for human action recognition. *IEEE transactions on pattern analysis and machine intelligence*, 35(1), 221-231.
- [Karpathy '14] Karpathy, A., Toderici, G., Shetty, S., Leung, T., Sukthankar, R., & Fei-Fei, L. (2014). Large-scale video classification with convolutional neural networks. In Proceedings of the IEEE conference on Computer Vision and Pattern Recognition (pp. 1725-1732).
- [Harris '88] Harris, C. G., & Stephens, M. (1988). A Combined Corner and Edge Detector. In *Alvey vision conference*. Vol. 15, No. 50, (pp. 147-151).
- [Matas '04] Matas, J., Chum, O., Urban, M., & Pajdla, T. (2004). Robust Wide-baseline Stereo from Maximally Stable Extremal Regions. *Image and Vision Computing*, 22(10), 761-767.
- [Dalal '05] Dalal, N., & Triggs, B. (2005). Histograms of Oriented Gradients for Human Detection. In *International Conference on Computer Vision & Pattern Recognition (CVPR'05)* (Vol. 1, pp. 886-893).
- [Lowe '99] Lowe, D. G. (1999). Object Recognition from Local Scale-Invariant Features. In *Proceedings of the International Conference on Computer Vision*, (Vol. 99, No. 2, pp. 1150-1157).
- [Kornilov '18] Kornilov, A. S., & Safonov, I. V. (2018). An overview of watershed algorithm implementations in open source libraries. *Journal of Imaging*, 4(10), 123.
- [Morlet '82] Morlet, J., Arens, G., Fourgeau, E., & Glard, D. (1982). Wave propagation and sampling theory—Part I: Complex signal and scattering in multilayered media. *Geophysics*, 47(2), 203-221.
- [Daubechies '92] Daubechies, I. (1992). Ten Lectures on Wavelets. *Society for Industrial and Applied Mathematics*.
- [Heil '89] Heil, Christopher E., and David F. Walnut. Continuous and discrete wavelet transforms. *SIAM review* 31.4 (1989): 628-666.
- [Haar '10] Haar, Alfréd (1910), Zur Theorie der orthogonalen Funktionensysteme, *Mathematische Annalen*, 69(3): 331–371.
- [Keogh '01] Keogh, E., Chakrabarti, K., Pazzani, M., & Mehrotra, S. (2001). Dimensionality reduction for fast similarity search in large time series databases. *Knowledge and Information Systems*, 3(3), 263-286.
- [Lin '03] Jessica, L., Keogh, E., Lonardi, S., and Chiu, B. (2003). A Symbolic Representation of Time Series, with Implications for Streaming Algorithms. In *Proceedings of the 8th ACM SIGMOD workshop on Research issues in data mining and knowledge discovery*, pp. 2-11
- [Wang '18] Wang, X., & Gupta, A. (2018). Videos as space-time region graphs. In Proceedings of the European conference on computer vision (ECCV) (pp. 399-417).
- [Wang '16] Wang, X., Gao, L., Song, J., & Shen, H. (2016). Beyond frame-level CNN: saliency-aware 3-D CNN with LSTM for video action recognition. *IEEE signal processing letters*, 24(4), 510-514.
- [Page '54] Page, E. S. (1954). Continuous Inspection Scheme. *Biometrika*. 41 (1/2): 100–115.
- [Aminikhanghahi '17] Aminikhanghahi, S., & Cook, D. J. (2017). A Survey of Methods for Time Series Change Point Detection. *Knowledge and information systems*, 51(2), 339-367.
- [Krishnamoorthy '13] Krishnamoorthy, A., & Menon, D. (2013). Matrix Inversion Using Cholesky Decomposition. In *2013 Signal Processing: Algorithms, Architectures, Arrangements, and Applications (SPA)* (pp. 70-72).

———— Chapter 7 ————

- [Bellman '61] Bellman, R. (1961). Curse of dimensionality. *Adaptive control processes: a guided tour*. Princeton University Press, 3, 2.
- [Deerwester '90] Deerwester, S., Dumais, S. T., Furnas, G. W., Landauer, T. K., & Harshman, R. (1990). Indexing by latent semantic analysis. *Journal of the American society for information science*, 41(6), 391-407.
- [Fisher '36] Fisher, R. (1936). Linear discriminant analysis. *Ann. Eugenics*, 7, 179.
- [Blei '03] Blei, D. M., Ng, A. Y., & Jordan, M. I. (2003). Latent Dirichlet Allocation. *Journal of machine Learning research*, 3(Jan), 993-1022
- [Cai '08] Cai, D., He, X., & Han, J. (2008). Training Linear Discriminant Analysis in Linear Time. In *2008 IEEE 24th International Conference on Data Engineering* (pp. 209-217).
- [Roweis '00] Roweis, S. T., & Saul, L. K. (2000). Nonlinear dimensionality reduction by locally linear embedding. *science*, 290 (5500), 2323-2326.
- [Smith '97] Smith, S. W. (1997) *The Scientists & Engineer's Guide to Digital Signal Processing*, California Technical Pub. San Diego..
- [Cholesky '10] Cholesky, A.L. (1910) On Numerical Solutions of Systems of Linear Equations. Available: <http://bibnum.education.fr/mathematiques/algebre/sur-la-resolution-numerique-des-systemes-d-equations-lineaires>
- [Sra '06] Sra, S., & Dhillon, I. S. (2006). Generalized Nonnegative Matrix Approximations with Bregman divergences. In *Advances in Neural Information Processing Systems* (pp. 283-290).
- [Stewart '93] Stewart, G. W. (1993). On the Early History of the Singular Value Decomposition. *SIAM review*, 35(4), 551-566.
- [Golub '71] Golub, G. H., & Reinsch, C. (1971). Singular Value Decomposition and Least Squares Solutions. In *Linear Algebra*(pp. 134-151). Springer, Berlin, Heidelberg.
- [Rehman '16] Ur Rehman, M. H., Liew, C. S., Abbas, A., Jayaraman, P. P., Wah, T. Y., & Khan, S. U. (2016). Big data reduction methods: a survey. *Data Science and Engineering*, 1(4), 265-284.
- [Kolda '09] Kolda, Tamara G., and Brett W. Bader. "Tensor decompositions and applications." *SIAM review* 51.3 (2009): 455-500.
- [Hitchcock '27] Hitchcock, L. (1927) The expression of a tensor or a polyadic as a sum of products, *Journal of Math and Physics*, 6, pp. 164–189.
- [Harshman '70] Harshman, R. A. (1970) Foundations of the PARAFAC procedure: Models and conditions for an "explanatory" multimodal factor analysis. *UCLA Working Papers in Phonetics*. pp. 1-84.
- [Carroll '70] Carroll, D. and Chang, J. J. (1970) Analysis of Individual Differences in Multidimensional Scaling via an N-way Generalization of "Eckart-Young" Decomposition, *Psychometrika*, 35, pp. 283–319.
- [Håstad '90] Håstad, J. (1990). Tensor rank is NP-complete. *Journal of Algorithms*, 11(4), 644-654.
- [Tucker '66] Tucker, L. R. (1966). Some mathematical notes on three-mode factor analysis. *Psychometrika*, 31(3), 279-311.
- [Kroonenberg '80] Kroonenberg, P. M., & De Leeuw, J. (1980). Principal Component Analysis of Three-mode Data by Means of Alternating Least Squares Algorithms. *Psychometrika*, 45(1), 69-97.
- [Oseledets '11] Oseledets, I. V. (2011). Tensor-Train Decomposition. *SIAM Journal on Scientific Computing*, 33(5), 2295-2317.

———— Chapter 8 ————

- [James '13] James, G., Witten, D., Hastie, T., & Tibshirani, R. (2013). *An Introduction to Statistical Learning: with Applications in R*. New York: springer.
- [Box '70] Box, G. E., Jenkins, G. M., Reinsel, G. C., & Ljung, G. M. (2015). *Time series analysis: forecasting and control*. John Wiley & Sons.
- [Dickey '79] Dickey, D. A., & Fuller, W. A. (1979). Distribution of the estimators for autoregressive time series with a unit root. *Journal of the American statistical association*, 74(366a), 427-431.
- [Domencich '75] Domencich, T. A., & McFadden, D. (1975). Urban travel demand-a behavioral analysis.
- [Wald '45] Wald, A. (1945). Sequential tests of statistical hypotheses. *The annals of mathematical statistics*, 16(2), 117-186.

- [Akaike '74] Akaike, H. (1974). A new look at statistical model identification. *IEEE Transactions on Automatic Control* 19, 716–723.
- [Schwarz '78] Schwarz, G. (1978). Estimating the dimension of a model. *The annals of statistics*, 6(2), 461–464.
- [Grus '19] Grus, J. (2019) Data Science from Scratch: First Principles with Python. *O'Reilly Media*.
- [Domingos '12] Domingos, P. M. (2012). A few useful things to know about machine learning. *Communication of ACM* 55(10), 78–87.
- [Thikhonov '43] Thikhonov, A. H. (1943). Об устойчивости обратных задач. In Докл. *Proceedings of the USSR Academy of Sciences* (Vol. 39, No. 5, pp. 195–198).
- [Santosa '86] Santosa, F., Symes, W.W. (1986). Linear Inversion of Band-Limited Reflection Seismograms". *SIAM Journal on Scientific and Statistical Computing*. SIAM. 7 (4): 1307–1330.
- [Tibshirani '96] Tibshirani, R. (1996). Regression Shrinkage and Selection via the LASSO. *Journal of the Royal Statistical Society. Series B (methodological)*. Wiley. 58 (1): 267–88.
- [Efron '04] Efron, B., Hastie, T., Johnstone, I., & Tibshirani, R. (2004). Least angle regression. *The Annals of statistics*, 32(2), 407–499.
- [Zou '05] Zou, H., & Hastie, T. (2005). Regularization and variable selection via the elastic net. *Journal of the royal statistical society: series B (statistical methodology)*, 67(2), 301–320.
- [Breiman '95] Breiman, L. (1995). Better Subset Regression using the Nonnegative Garrote. *Technometrics*, 37(4), 373–384.
- [Xiong '10] Xiong, S. (2010). Some Notes on the Nonnegative Garrote. *Technometrics*, 52(3), 349–361.
- [Robbins '51] Robbins, H., & Monro, S. (1951). A stochastic approximation method. *The annals of mathematical statistics*, 400–407.
- [Burkov '19] Burkov, A. (2019). The Hundred-Page Machine Learning Book. Andriy Burkov.
- [Kochenderfer '19] Kochenderfer, M. J., & Wheeler, T. A. (2019). *Algorithms for optimization*. MIT Press.
- [Kingma '14] Kingma, D. P., & Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.
- [Brownlee '16] Brownlee, J. (2016). Master Machine Learning Algorithms: discover how they work and implement them from scratch. Machine Learning Mastery.

———— Chapter 9 ————

- [Bayes '63] Bayes, T. (1763). An essay towards solving a problem in the doctrine of chances. *Philosophical Transactions of the Royal Society of London*, 53, 370–418.
- [Cover '67] Cover, T., & Hart, P. (1967). Nearest neighbor pattern classification. *IEEE Transactions on Information Theory*, 13(1), 21–27.
- [Voronoi '08] Voronoi, G. (1908). Nouvelles applications des paramètres continus à la théorie des formes quadratiques. Deuxième mémoire. Recherches sur les paralléloèdres primitifs. *Journal für die Reine und Angewandte Mathematik*, 134, 198–287.
- [Bentley '75] Bentley, J. L. (1975). Multidimensional binary search trees used for associative searching. *Communications of the ACM*, 18(9), 509–517.
- [Indyk '98] Indyk, P., & Motwani, R. (1998). *Approximate nearest neighbors: Towards removing the curse of dimensionality*. Proceedings of the Thirtieth Annual ACM Symposium on Theory of Computing (STOC '98), 604–613.
- [Cortes '95] Cortes, C., & Vapnik, V. (1995). Support-Vector Networks. *Machine learning*, 20(3), 273–297.
- [Bordes '05] Bordes, A., Ertekin, S., Weston, J., & Bottou, L. (2005). Fast kernel classifiers with online and active learning. *Journal of Machine Learning Research*, 6(Sep), 1579–1619.
- [Scholkopf '01] Scholkopf, B., & Smola, A. J. (2001). Learning with kernels: support vector machines, regularization, optimization, and beyond. MIT press.
- [Kass '80] Kass, Gordon V.; An Exploratory Technique for Investigating Large Quantities of Categorical Data, *Applied Statistics*, Vol. 29, No. 2 (1980), pp. 119–127
- [Quinlan '94] Quinlan, J. R. (1994). C4. 5: Programs for Machine Learning. *Elsevier*

- [Breimanetal '84] Breimanetal, L. (1984). Classification and Regression Trees. Wadsworth, Belmont.
- [Breiman '96] Breiman, L. (1996). Bagging predictors. *Machine Learning*, 24(2), 123-140.
- [Schapire '90] Schapire, R. E. (1990). The strength of weak learnability. *Machine Learning*, 5(2), 197-227.
- [Wolpert '92] Wolpert, D. H. (1992). Stacked generalization. *Neural networks*, 5(2), 241-259.
- [Breiman '05] Breiman, L. (2001). Random Forests. *Machine learning*, 45(1), 5-32.
- [Freund '95] Freund, Y., & Schapire, R. E. (1995). A Desicion-Theoretic Generalization of On-line Learning and an Application to Boosting. In *European conference on computational learning theory* (pp. 23-37). Springer, Berlin, Heidelberg.
- [Friedman '01] Friedman, J. H. (2001). Greedy Function Approximation: A Gradient Boosting Machine. *Annals of statistics*, 1189-1232.
- [Si '17] Si, S., Zhang, H., Keerthi, S. S., Mahajan, D., Dhillon, I. S., & Hsieh, C. J. (2017). Gradient boosted decision trees for high dimensional sparse output. In International conference on machine learning (pp. 3182-3190). PMLR.
- [Mason '99] Mason, L., Baxter, J., Bartlett, P., & Frean, M. (1999, May). Boosting algorithms as gradient descent in function space. NIPS.
- [Chen '15] Chen, T., He, T., Benesty, M., Khotilovich, V., Tang, Y., & Cho, H. (2015). XGBoost: extreme gradient boosting. *R package version 0.4-2*, 1(4).
- [Ke '17] Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., ... & Liu, T. Y. (2017). LightGBM: A highly efficient gradient boosting decision tree. *Advances in neural information processing systems*, 30, 3146-3154.
- [Prokhorenkova '18] Prokhorenkova, L., Gusev, G., Vorobev, A., Dorogush, A. V., & Gulin, A. (2018). CatBoost: unbiased boosting with categorical features. *Advances in Neural Information Processing Systems*, 31.
- [Boehmke '20] Boehmke, B., & Greenwell, B. M. (2019). Hands-on machine learning with R. CRC Press.
- [Keshavarz '21] Keshavarz, H., Abadeh, M. S., & Rawassizadeh, R. (2021). SEFR: A fast linear-time classifier for ultra-low power devices. *arXiv preprint arXiv:2006.04620*.

———— Chapter 10 ————

- [McCulloch '43] McCulloch, W. S., & Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. *The bulletin of mathematical biophysics*, 5(4), 115-133.
- [Rosenblatt '58] Rosenblatt, F. (1958). The perceptron: a probabilistic model for information storage and organization in the brain. *Psychological review*, 65(6), 386.
- [Minsky '69] Minsky, M., & Papert, S. A. (1969). Perceptrons: An introduction to computational geometry. MIT press.
- [Ciresan '11] Ciresan, D. C., Meier, U., Masci, J., Gambardella, L. M., Schmidhuber, J. Flexible, High Performance Convolutional Neural Networks for Image Classification. (2011) *International Joint Conference on Artificial Intelligence (IJCAI-2011)*
- [Krizhevsky '12] Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. In *Advances in neural information processing systems* (pp. 1097-1105).
- [Aggarwal '18] Aggarwal, C. C. (2018). Neural networks and deep learning. *Springer*.
- [Hofstadter '79] Hofstadter, D. R. (1979). Gödel, Escher, Bach: an eternal golden braid (Vol. 13). New York: Basic books.
- [Rawassizadeh '19] Rawassizadeh, R., Sen, T., Kim, S. J., Meurisch, C., Keshavarz, H., Mühlhäuser, M., & Pazzani, M. (2019). Manifestation of virtual assistants and robots into daily life: Vision and challenges. *CCF Transactions on Pervasive Computing and Interaction*, 1-12.
- [Rasheed '16] Rasheed, T. (2016) Make Your Own Neural Network. *CreateSpace Independent Publishing Platform*.
- [Chollet '18] Chollet, F. (2017). Deep Learning with Python, *Manning Publications CO*.
- [Nair '10] Nair, V., & Hinton, G. E. (2010). Rectified Linear Units Improve Restricted Boltzmann Machines. In *ICML*.
- [Maas '13] Maas, A. L., Hannun, A. Y., & Ng, A. Y. (2013). Rectifier Nonlinearities Improve Neural Network Acoustic Models. In *Proc. icml* (Vol. 30, No. 1, p. 3).

- [Hellinger '09] Hellinger, E. (1909), Neue Begründung der Theorie quadratischer Formen von unendlichvielen Veränderlichen, *Journal für die reine und angewandte Mathematik* (in German), 136: 210–271
- [Bharath '99] Bharath, B., & Borkar, V. S. (1999). Stochastic Approximation Algorithms: Overview and Recent Trends. *Sadhana*, 24(4-5), 425-452.
- [Polyak '64] Polyak, B. T. (1964). Some Methods of Speeding up the Convergence of Iteration Methods. *USSR Computational Mathematics and Mathematical Physics*, 4(5), 1-17.
- [Nesterov '83] Nesterov, Y. A Method of Solving a Convex Programming Problem with Convergence Rate $O(1/\sqrt{k})$. *Soviet Mathematics Doklady*, 27:372–376, 1983.
- [Sutskever '13] Sutskever, I., Martens, J., Dahl, G., & Hinton, G. (2013). On the Importance of Initialization and Momentum in Deep Learning. In *International Conference on Machine Learning* (pp. 1139-1147).
- [Duchi '11] Duchi, J., Hazan, E., & Singer, Y. (2011). Adaptive Subgradient Methods for Online Learning and Stochastic Optimization. *Journal of machine learning research*, 12(7).
- [Zeiler '12] Zeiler, M. D. (2012). Adadelta: an Adaptive Learning Rate Method. *arXiv preprint arXiv:1212.5701*.
- [Hinton '12] Hinton, G., Srivastava, N., & Swersky, K. (2012). RMSprop: Divide the Gradient by a Running Average of its Recent Magnitude. *Neural Networks for Machine Learning, Lecture 6e*.
- [Riedmiller '93] Riedmiller, M., & Braun, H. (1993). A Direct Adaptive Method for Faster Backpropagation Learning: The RPROP algorithm. In *IEEE International Conference on Neural Networks* (pp. 586-591).
- [Schmidt '20] Schmidt, R. M., Schneider, F., & Hennig, P. (2020). Descending through a Crowded Valley-- Benchmarking Deep Learning Optimizers. *arXiv preprint arXiv:2007.01547*.
- [Chen '19] Chen, J., Kyrillidis, A. Decaying Momentum Helps Neural Network Training. *arXiv preprint arXiv:1910.04952*, 2019.
- [Zhang '17] Zhang, J., & Mitliagkas, I. (2017). Yellowfin and the Art of Momentum Tuning. *arXiv preprint arXiv:1706.03471*.
- [Schneider '19] Schneider, Frank, Lukas Balles, and Philipp Hennig. (2019) DeepOBS: A Deep Learning Optimizer Benchmark Suite. *International Conference on Learning Representations*.
- [Kelley '60] Kelley, H. J. (1960). Gradient theory of optimal flight paths. *Ars Journal*, 30(10), 947-954.
- [Bryson '62] Bryson, A. E. (1961, April). A Gradient Method for Optimizing Multi-stage Allocation Processes. In *Proc. Harvard Univ. Symposium on digital computers and their applications* (Vol. 72).
- [Rumelhart '86] Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986). Learning Representations by Back-Propagating Errors. *nature*, 323(6088), 533-536.
- [Glorot '10] Glorot, X., & Bengio, Y. (2010, March). Understanding the Difficulty of Training Deep Feedforward Neural Networks. In *Proceedings of the thirteenth international conference on artificial intelligence and statistics* (pp. 249-256).
- [He '15] He, K., Zhang, X., Ren, S., & Sun, J. (2015). Delving Deep into Rectifiers: Surpassing Human-Level Performance on Imagenet Classification. In *Proceedings of the IEEE International Conference on Computer Vision* (pp. 1026-1034).
- [Srivastava '14] Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. (2014). Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *The journal of machine learning research*, 15(1), 1929-1958.
- [Gal '16] Gal, Y., & Ghahramani, Z. (2016). Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning. In *International Conference on Machine Learning* (pp. 1050-1059).
- [Konečný '16] Konečný, J., McMahan, H. B., Yu, F. X., Richtárik, P., Suresh, A. T., & Bacon, D. (2016). Federated learning: Strategies for improving communication efficiency. *arXiv preprint arXiv:1610.05492*.
- [Warden '19] Warden, P., & Situnayake, D. (2019). TinyML: Machine learning with tensorflow lite on Arduino and ultra-low-power microcontrollers. O'Reilly Media, Inc.
- [Hubel '59] Hubel, DH; Wiesel, T.N. (1959). Receptive Fields of single Neurones in the Cat's Striate Cortex. *J. Physiol.* 148 (3): 574–91.
- [Fukushima '80] Fukushima, K. (1980). "Neocognitron: A Self-organizing Neural Network Model for a Mechanism of Pattern Recognition Unaffected by Shift in Position" (PDF). *Biological Cybernetics*. **36** (4): 193–202.

- [LeCun '89] Y. LeCun, B. Boser, J. S. Denker, D. Henderson, R. E. Howard, W. Hubbard, L. D. Jackel (1989) Backpropagation Applied to Handwritten Zip Code Recognition.
- [Deng '09] Deng, J., Dong, W., Socher, R., Li, L. J., Li, K., & Fei-Fei, L. (2009, June). Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on Computer Vision and Pattern Recognition* (pp. 248-255).
- [Brown '20] Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., ... & Agarwal, S. (2020). Language models are few-shot learners. *arXiv preprint arXiv:2005.14165*.
- [Hochreiter '97] Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural computation*, 9(8), 1735-1780.
- [Greff '16] Greff, K., Srivastava, R. K., Koutník, J., Steunebrink, B. R., & Schmidhuber, J. (2016). LSTM: A search space odyssey. *IEEE transactions on neural networks and learning systems*, 28(10), 2222-2232.
- [Chung '14] Chung, J., Gulcehre, C., Cho, K., & Bengio, Y. (2014). Empirical evaluation of gated recurrent neural networks on sequence modeling. *arXiv preprint arXiv:1412.3555*.
- [Karpathy '15] Karpathy, A. (2015). The Unreasonable Effectiveness of Recurrent Neural Networks. *Andrej Karpathy blog*, 21, 23.
- [Dumoulin '16] Dumoulin, V., & Visin, F. (2016). A guide to convolution arithmetic for deep learning. *arXiv preprint arXiv:1603.07285*.
- [Chen '17] Chen, L. C., Papandreou, G., Schroff, F., & Adam, H. (2017). Rethinking atrous convolution for semantic image segmentation. *arXiv preprint arXiv:1706.05587*.
- [Ramesh '22] Ramesh, A., Dhariwal, P., Nichol, A., Chu, C., & Chen, M. (2022). Hierarchical text-conditional image generation with clip latents. *arXiv preprint arXiv:2204.06125*.

———— Chapter 11 ————

- [Foster '19] Foster, D. (2019). Generative deep learning: teaching machines to paint, write, compose, and play. O'Reilly Media.
- [Kohonen '82] Kohonen, T. (1982). Self-organized formation of topologically correct feature maps. *Biological cybernetics*, 43(1), 59-69.
- [Tian '14] Tian, J., Azarian, M. H., & Pecht, M. (2014). Anomaly detection using self-organizing maps-based k-nearest neighbor algorithm. In *Proceedings of the European Conference of the Prognostics and Health Management Society* (pp. 1-9). Citeseer.
- [LeCun '06] LeCun, Y., Chopra, S., Hadsell, R., Ranzato, M., & Huang, F. (2006). A tutorial on energy-based learning. *Predicting structured data*, 1(0).
- [Ackley '85] Ackley, D. H., Hinton, G. E., & Sejnowski, T. J. (1985). A learning algorithm for Boltzmann machines. *Cognitive science*, 9(1), 147-169.
- [Gibbs '02] Gibbs, J. W. (1902). *Elementary principles in statistical mechanics*: developed with especial reference to the rational foundations of thermodynamics. C. Scribner's Sons.
- [Hinton '02] Hinton, G. E. (2002). Training products of experts by minimizing contrastive divergence. *Neural computation*, 14(8), 1771-1800.
- [Hinton '95] Hinton, G. E., Dayan, P., Frey, B. J., & Neal, R. M. (1995). The "wake-sleep" algorithm for unsupervised neural networks. *Science*, 268(5214), 1158-1161.
- [Wang '17] Wang, H., & Raj, B. (2017). On the Origin of Deep Learning. *arXiv preprint arXiv:1702.07800*.
- [Hopfield '82] Hopfield, J. J. (1982). Neural Networks and Physical Systems with Emergent Collective Computational Abilities. *Proceedings of the national academy of sciences*, 79(8), 2554-2558.
- [Bengio '09] Bengio, Y. (2009). Learning Deep Architectures for AI. Now Publishers Inc.
- [Kingma '13] Kingma, D. P., & Welling, M. (2013). Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*.
- [Makhzani '13] Makhzani, A., & Frey, B. (2013). K-sparse autoencoders. *arXiv preprint arXiv:1312.5663*.
- [Ng '11] Ng, A. (2011). Sparse autoencoder. CS294A Lecture notes, 72 (2011), 1-19.
- [Vincent '10] Vincent, P., Larochelle, H., Lajoie, I., Bengio, Y., Manzagol, P. A., & Bottou, L. (2010). Stacked denoising autoencoders: Learning useful representations in a deep network with a local denoising criterion. *Journal of machine learning research*, 11(12).

- [Rifai '11] Rifai, S., Vincent, P., Muller, X., Glorot, X., & Bengio, Y. (2011). Contractive auto-encoders: Explicit invariance during feature extraction. In *International Conference of Machine Learning*.
- [Bengio '07] Bengio, Y., Lamblin, P., Popovici, D., & Larochelle, H. (2006). Greedy layer-wise training of deep networks. *Advances in neural information processing systems*, 19.
- [Hinton '06] Hinton, G. E., & Salakhutdinov, R. R. (2006). Reducing the dimensionality of data with neural networks. *science*, 313(5786), 504-507.
- [Foster '19] Foster, D. (2019). *Generative Deep Learning: Teaching Machines to Paint, Write, Compose, and Play*. O'Reilly Media.
- [Langr '19] Langr, J., & Bok, V. (2019). *GANs in action: deep learning with generative adversarial networks*. Manning Publications.
- [Ramesh '21] Ramesh, A., Pavlov, M., Goh, G., Gray, S., Voss, C., Radford, A., ... & Sutskever, I. (2021). Zero-shot text-to-image generation. arXiv preprint arXiv:2102.12092.
- [Ronneberger '15] Ronneberger, O., Fischer, P., & Brox, T. (2015). U-net: Convolutional networks for biomedical image segmentation. In *International Conference on Medical image computing and computer-assisted intervention* (pp. 234-241).
- [Goodfellow '14] Goodfellow, I. J., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., & Bengio, Y. (2014). Generative adversarial networks. arXiv preprint arXiv:1406.2661.
- [Niemitalo '11] Niemitalo, O. (2010). A method for training artificial neural networks to generate missing data within a variable context. *Blog post, Internet Archive*. <https://web.archive.org/web/20120312111546/http://yehar.com:80/blog/?p=167>
- [Li '13] Li, W., Gauci, M., & Groß, R. (2013). A coevolutionary approach to learn animal behavior through controlled interaction. In *Proceedings of the 15th annual conference on Genetic and Evolutionary Computation* (pp. 223-230).
- [Karras '19] Karras, T., Laine, S., & Aila, T. (2019). A style-based generator architecture for generative adversarial networks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 4401-4410).
- [Salimans '16] Salimans, T., Goodfellow, I., Zaremba, W., Cheung, V., Radford, A., & Chen, X. (2016). Improved techniques for training GANs. arXiv preprint arXiv:1606.03498.
- [Heusel '17] Heusel, M., Ramsauer, H., Unterthiner, T., Nessler, B., & Hochreiter, S. (2017). Gans trained by a two time-scale update rule converge to a local nash equilibrium. arXiv preprint arXiv:1706.08500.
- [Szegedy '16] Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., & Wojna, Z. (2016). Rethinking the inception architecture for computer vision. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 2818-2826).
- [Mirza '14] Mirza, M., & Osindero, S. (2014). Conditional generative adversarial nets. arXiv preprint arXiv:1411.1784.
- [Rüschendorf '85] Rüschendorf, L. (1985). The Wasserstein distance and approximation theorems. *Probability Theory and Related Fields*, 70(1), 117-129.
- [Gulrajani '17] Gulrajani, I., Ahmed, F., Arjovsky, M., Dumoulin, V., & Courville, A. (2017). Improved training of wasserstein gans. arXiv preprint arXiv:1704.00028.
- [Isola '17] Isola, P., Zhu, J. Y., Zhou, T., & Efros, A. A. (2017). Image-to-image translation with conditional adversarial networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* (pp. 1125-1134).
- [Zhu '17] Zhu, J. Y., Park, T., Isola, P., & Efros, A. A. (2017). Unpaired image-to-image translation using cycle-consistent adversarial networks. In *Proceedings of the IEEE international conference on computer vision* (pp. 2223-2232).
- [Karras '20] Karras, T., Laine, S., Aittala, M., Hellsten, J., Lehtinen, J., & Aila, T. (2020). Analyzing and improving the image quality of stylegan. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 8110-8119).
- [Huang '17] Huang, X., & Belongie, S. (2017). Arbitrary style transfer in real-time with adaptive instance normalization. In *Proceedings of the IEEE International Conference on Computer Vision* (pp. 1501-1510).

- [Karras '17] Karras, T., Aila, T., Laine, S., & Lehtinen, J. (2017). Progressive growing of gans for improved quality, stability, and variation. *arXiv preprint arXiv:1710.10196*.
- [Karras '21] Karras, T., Aittala, M., Laine, S., Härkönen, E., Hellsten, J., Lehtinen, J., & Aila, T. (2021). Alias-Free Generative Adversarial Networks. *arXiv preprint arXiv:2106.12423*.
- [Brownlee '19] Brownlee, J. (2019). Generative Adversarial Networks with Python: Deep Learning Generative Models for Image Synthesis and Image Translation. Machine Learning Mastery.
- [Zakharov '19] Zakharov, E., Shysheya, A., Burkov, E., & Lempitsky, V. (2019). Few-shot adversarial learning of realistic neural talking head models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* (pp. 9459-9468).
- [Bengio '13] Bengio, Y., Courville, A., & Vincent, P. (2013). Representation learning: A review and new perspectives. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 35(8), 1798-1828.
- [Le-Khac '20] Le-Khac, P. H., Healy, G., & Smeaton, A. F. (2020). Contrastive representation learning: A framework and review. *IEEE Access*, 8, 193907-193934.
- [Bell '10] Bell, G., & Gemmell, J. (2010). Your life, uploaded: The digital way to better memory, health, and productivity. Penguin.
- [Chen '20 A] Chen, T., Kornblith, S., Norouzi, M., & Hinton, G. (2020). A simple framework for contrastive learning of visual representations. In *International conference on machine learning* (pp. 1597-1607). PMLR.
- [Chen '20 B] Chen, T., Kornblith, S., Swersky, K., Norouzi, M., & Hinton, G. E. (2020). Big self-supervised models are strong semi-supervised learners. *Advances in neural information processing systems*, 33, 22243-22255.
- [He '20] He, K., Fan, H., Wu, Y., Xie, S., & Girshick, R. (2020). Momentum contrast for unsupervised visual representation learning. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition* (pp. 9729-9738).
- [Chen '20 C] Chen, X., Fan, H., Girshick, R., & He, K. (2020). Improved baselines with momentum contrastive learning. *arXiv preprint arXiv:2003.04297*.
- [Chopra '05] Chopra, S., Hadsell, R., & LeCun, Y. (2005). Learning a Similarity Metric Discriminatively, with Application to Face Verification. In *2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05)* (Vol. 1, pp. 539-546).
- [Schroff '15] Schroff, F., Kalenichenko, D., & Philbin, J. (2015). Facenet: A unified embedding for face recognition and clustering. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 815-823).
- [Radford '21] Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... & Sutskever, I. (2021). Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning* (pp. 8748-8763). PMLR.
- [Crowson '22] Crowson, K., Biderman, S., Kornis, D., Stander, D., Hallahan, E., Castricato, L., & Raff, E. (2022). Vqgan-clip: Open domain image generation and editing with natural language guidance. *arXiv preprint arXiv:2204.08583*.
- [Van Den Oord '17] Van Den Oord, A., & Vinyals, O. (2017). Neural discrete representation learning. *Advances in neural information processing systems*, 30.
- [Esser '21] Esser, P., Rombach, R., & Ommer, B. (2021). Taming transformers for high-resolution image synthesis. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 12873-12883).
- [Ramesh '21] Ramesh, Aditya, Mikhail Pavlov, Gabriel Goh, Scott Gray, Chelsea Voss, Alec Radford, Mark Chen, and Ilya Sutskever. (2021) "Zero-shot text-to-image generation." In *International Conference on Machine Learning*, pp. 8821-8831. PMLR.
- [Maddison '16] Maddison, C. J., Mnih, A., & Teh, Y. W. (2016). The concrete distribution: A continuous relaxation of discrete random variables. *arXiv preprint arXiv:1611.00712*.
- [Jang '16] Jang, E., Gu, S., & Poole, B. (2016). Categorical reparameterization with gumbel-softmax. *arXiv preprint arXiv:1611.01144*.
- [Bishop '24] Bishop, C. M., & Bishop, H. (2024). Deep learning: foundations and concepts. *Springer*.
- [Ramesh '22] Ramesh, A., Dhariwal, P., Nichol, A., Chu, C., & Chen, M. (2022). Hierarchical text-conditional image generation with clip latents. *arXiv preprint arXiv:2204.06125*.

- [Nichol '21] Nichol, A., Dhariwal, P., Ramesh, A., Shyam, P., Mishkin, P., McGrew, B., ... & Chen, M. (2021). Glide: Towards photorealistic image generation and editing with text-guided diffusion models. arXiv preprint arXiv:2112.10741.
- [Rombach '22] Rombach, R., Blattmann, A., Lorenz, D., Esser, P., & Ommer, B. (2022). High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 10684-10695).
- [Zhang '21] Zhang, K., Liang, J., Van Gool, L., & Timofte, R. (2021). Designing a practical degradation model for deep blind image super-resolution. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* (pp. 4791-4800).
- [Saharia '22] Saharia, C., Chan, W., Saxena, S., Li, L., Whang, J., Denton, E., ... & Norouzi, M. (2022). Photorealistic Text-to-Image Diffusion Models with Deep Language Understanding. arXiv preprint arXiv:2205.11487.
- [Yu '22] Yu, J., Xu, Y., Koh, J. Y., Luong, T., Baid, G., Wang, Z., ... & Wu, Y. (2022). Scaling Autoregressive Models for Content-Rich Text-to-Image Generation. arXiv preprint arXiv:2206.10789.
- [Dhariwal '21] Dhariwal, P., & Nichol, A. (2021). Diffusion models beat GANs on image synthesis. *Advances in Neural Information Processing Systems*, 34, 8780-8794.
- [Yu '21] Yu, J., Li, X., Koh, J. Y., Zhang, H., Pang, R., Qin, J., ... & Wu, Y. (2021). Vector-quantized image modeling with improved VQGAN. arXiv preprint arXiv:2110.04627.
- [Betker '23] Betker, J., Goh, G., Jing, L., Brooks, T., Wang, J., Li, L., ... & Ramesh, A. (2023). Improving image generation with better captions. Computer Science. <https://cdn.openai.com/papers/dall-e-3.pdf>, 2(3), 8.
- [Rombach '22] Rombach, R., Blattmann, A., Lorenz, D., Esser, P., & Ommer, B. (2022). High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition* (pp. 10684-10695).
- [Podell '23] Podell, D., English, Z., Lacey, K., Blattmann, A., Dockhorn, T., Müller, J., ... & Rombach, R. (2023). Sdxl: Improving latent diffusion models for high-resolution image synthesis. arXiv preprint arXiv:2307.01952.
- [Zhu '23] Zhu, Y., Baca, J., Rekabdar, B., & Rawassizadeh, R. (2023). A Survey of AI Music Generation Tools and Models. arXiv preprint arXiv:2308.12982.
- [Singer '22] Singer, U., Polyak, A., Hayes, T., Yin, X., An, J., Zhang, S., ... & Taigman, Y. (2022). Make-a-video: Text-to-video generation without text-video data. arXiv preprint arXiv:2209.14792.
- [Khachatryan '23] Khachatryan, L., Movsisyan, A., Tadevosyan, V., Henschel, R., Wang, Z., Navasardyan, S., & Shi, H. (2023). Text2video-zero: Text-to-image diffusion models are zero-shot video generators. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* (pp. 15954-15964).
- [Liu '24] Liu, Y., Zhang, K., Li, Y., Yan, Z., Gao, C., Chen, R., ... & Sun, L. (2024). Sora: A Review on Background, Technology, Limitations, and Opportunities of Large Vision Models. arXiv preprint arXiv:2402.17177.
- [Kandala '24] Kandala, H., Gao, J., & Yang, J. (2024). Pix2Gif: Motion-Guided Diffusion for GIF Generation. arXiv preprint arXiv:2403.04634.
- [Esser '24] Esser, P., Kulal, S., Blattmann, A., Entezari, R., Müller, J., Saini, H., ... & Rombach, R. (2024). Scaling rectified flow transformers for high-resolution image synthesis. arXiv preprint arXiv:2403.03206.
- [Sauer '24] Sauer, A., Boesel, F., Dockhorn, T., Blattmann, A., Esser, P., & Rombach, R. (2024). Fast High-Resolution Image Synthesis with Latent Adversarial Diffusion Distillation. arXiv preprint arXiv:2403.12015.

———— Chapter 12 ————

- [Sutskever '14] Sutskever, I., Vinyals, O., & Le, Q. V. (2014). Sequence to sequence learning with neural networks. In *Advances in neural information processing systems* (pp. 3104-3112).
- [Cho '14] Cho, K., Van Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H., & Bengio, Y. (2014). Learning phrase representations using RNN encoder-decoder for statistical machine translation. arXiv preprint arXiv:1406.1078.
- [Bahdanau '14] Bahdanau, D., Cho, K., & Bengio, Y. (2014). Neural machine translation by jointly learning to align and translate. arXiv preprint arXiv:1409.0473.
- [Luong '15] Luong, M. T., Pham, H., & Manning, C. D. (2015). Effective approaches to attention-based neural machine translation. arXiv preprint arXiv:1508.04025.

- [Lin '17] Lin, Z., Feng, M., Santos, C. N. D., Yu, M., Xiang, B., Zhou, B., & Bengio, Y. (2017). A structured self-attentive sentence embedding. arXiv preprint arXiv:1703.03130.
- [Vaswani '17] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. In *Advances in neural information processing systems* (pp. 5998-6008).
- [Peters '18] Peters, M. E., Neumann, M., Iyyer, M., Gardner, M., Clark, C., Lee, K., & Zettlemoyer, L. (2018). Deep contextualized word representations. arXiv preprint arXiv:1802.05365.
- [Howard '18] Howard, J., & Ruder, S. (2018). Universal language model fine-tuning for text classification. arXiv preprint arXiv:1801.06146.
- [Lewis '19] Lewis, M., Liu, Y., Goyal, N., Ghazvininejad, M., Mohamed, A., Levy, O., ... & Zettlemoyer, L. (2019). Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. arXiv preprint arXiv:1910.13461.
- [Gu '23] Gu, A., & Dao, T. (2023). Mamba: Linear-time sequence modeling with selective state spaces. arXiv preprint arXiv:2312.00752.
- [Ramachandran '17] Ramachandran, P., Zoph, B., & Le, Q. V. (2017). Swish: a self-gated activation function. arXiv preprint arXiv:1710.05941, 7(1), 5.
- [Hendrycks '16] Hendrycks, D., & Gimpel, K. (2016). Gaussian Error Linear Units (GeLUs). arXiv preprint arXiv:1606.08415.
- [Li '23] Li, P., Yang, J., Islam, M. A., & Ren, S. (2023). Making ai less" thirsty": Uncovering and addressing the secret water footprint of AI models. arXiv preprint arXiv:2304.03271.
- [Liu '19] Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., ... & Stoyanov, V. (2019). RoBERTa: A robustly optimized bert pretraining approach. arXiv preprint arXiv:1907.11692.
- [Lee '20] Lee, J., Yoon, W., Kim, S., Kim, D., Kim, S., So, C. H., & Kang, J. (2020). BioBERT: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4), 1234-1240.
- [Lan '19] Lan, Z., Chen, M., Goodman, S., Gimpel, K., Sharma, P., & Soricut, R. (2019). Albert: A lite bert for self-supervised learning of language representations. arXiv preprint arXiv:1909.11942.
- [Sanh '19] Sanh, V., Debut, L., Chaumond, J., & Wolf, T. (2019). DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. arXiv preprint arXiv:1910.01108.
- [Lample '19] Lample, G., & Conneau, A. (2019). Cross-lingual language model pretraining. arXiv preprint arXiv:1901.07291.
- [Raffel '20] Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., ... & Liu, P. J. (2020). Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer. *Journal of Machine Learning Research*, 21, 1-67.
- [Radford '18] Radford, A., Narasimhan, K., Salimans, T., & Sutskever, I. (2018). Improving language understanding by generative pre-training.
- [Zhu '15] Zhu, Y., Kiros, R., Zemel, R., Salakhutdinov, R., Urtasun, R., Torralba, A., & Fidler, S. (2015). Aligning books and movies: Towards story-like visual explanations by watching movies and reading books. In *Proceedings of the IEEE international conference on computer vision* (pp. 19-27).
- [Liu '18] Liu, P. J., Saleh, M., Pot, E., Goodrich, B., Sepassi, R., Kaiser, L., & Shazeer, N. (2018). Generating wikipedia by summarizing long sequences. arXiv preprint arXiv:1801.10198.
- [Radford '19] Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., & Sutskever, I. (2019). Language models are unsupervised multitask learners. *OpenAI blog*, 1(8), 9.
- [Larochelle '08] Larochelle, H., Erhan, D., & Bengio, Y. (2008). Zero-data learning of new tasks. In *AAAI* (Vol. 1, No. 2, p. 3).
- [Touvron '23] Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M. A., Lacroix, T., ... & Lample, G. (2023). Llama: Open and efficient foundation language models. arXiv preprint arXiv:2302.13971.
- [Hoffmann '22] Hoffmann, J., Borgeaud, S., Mensch, A., Buchatskaya, E., Cai, T., Rutherford, E., ... & Sifre, L. (2022). Training compute-optimal large language models. arXiv preprint arXiv:2203.15556.
- [Gao '20] Gao, L., Biderman, S., Black, S., Golding, L., Hoppe, T., Foster, C., ... & Leahy, C. (2020). The pile: An 800gb dataset of diverse text for language modeling. arXiv preprint arXiv:2101.00027.

- [Su '24] Su, J., Ahmed, M., Lu, Y., Pan, S., Bo, W., & Liu, Y. (2024). RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568, 127063.
- [Shazeer '20] Shazeer, N. (2020). Glu variants improve transformer. *arXiv preprint arXiv:2002.05202*. Chicago
- [Touvron '23] Touvron, H., Martin, L., Stone, K., Albert, P., Almahairi, A., Babaei, Y., ... & Scialom, T. (2023). Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*.
- [Dubey '24] Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., ... & Ganapathy, R. (2024). The Llama 3 Herd of Models. *arXiv preprint arXiv:2407.21783*.
- [Taori '23] Taori, R., Gulrajani, I., Zhang, T., Dubois, Y., Li, X., Guestrin, C., ... & Hashimoto, T. B. (2023). Alpaca: A strong, replicable instruction-following model. Stanford Center for Research on Foundation Models. <https://crfm.stanford.edu/2023/03/13/alpaca.html>, 3(6), 7.
- [Jiang '23] Jiang, A. Q., Sablayrolles, A., Mensch, A., Bamford, C., Chaplot, D. S., Casas, D. D. L., ... & Sayed, W. E. (2023). Mistral 7B. *arXiv preprint arXiv:2310.06825*.
- [Jiang '24] Jiang, A. Q., Sablayrolles, A., Roux, A., Mensch, A., Savary, B., Bamford, C., ... & Sayed, W. E. (2024). Mixtral of experts. *arXiv preprint arXiv:2401.04088*.
- [Signal '23] Signal, K., Azizi, S., Tu, T., Mahdavi, S. S., Wei, J., Chung, H. W., ... & Natarajan, V. (2023). Large language models encode clinical knowledge. *Nature*, 620(7972), 172-180.
- [Chang '08] Chang, M. W., Ratnoff, L. A., Roth, D., & Srikumar, V. (2008). Importance of Semantic Representation: Dataless Classification. In *Aaai* (Vol. 2, pp. 830-835).
- [Brown '20] Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., ... & Amodei, D. (2020). Language models are few-shot learners. *arXiv preprint arXiv:2005.14165*.
- [Child '19] Child, R., Gray, S., Radford, A., & Sutskever, I. (2019). Generating long sequences with sparse transformers. *arXiv preprint arXiv:1904.10509*.
- [LeCun '89] LeCun, Y., Boser, B., Denker, J. S., Henderson, D., Howard, R. E., Hubbard, W., & Jackel, L. D. (1989). Backpropagation applied to handwritten zip code recognition. *Neural computation*, 1(4), 541-551.
- [Devlin '18] Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2018). Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*.
- [Rafailov '24] Rafailov, R., Sharma, A., Mitchell, E., Manning, C. D., Ermon, S., & Finn, C. (2024). Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36.
- [Shafahi '19] Shafahi, A., Najibi, M., Ghiasi, M. A., Xu, Z., Dickerson, J., Studer, C., ... & Goldstein, T. (2019). Adversarial training for free!. *Advances in neural information processing systems*, 32.
- [Jelinek '77] Jelinek, F., Mercer, R. L., Bahl, L. R., & Baker, J. K. (1977). Perplexity—a measure of the difficulty of speech recognition tasks. *The Journal of the Acoustical Society of America*, 62(S1), S63-S63.
- [Morris '04] Morris, A. C., Maier, V., & Green, P. (2004). From WER and RIL to MER and WIL: improved evaluation measures for connected speech recognition. In *Eighth International Conference on Spoken Language Processing*.
- [Lin '04] Lin, C. Y. (2004). Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out* (pp. 74-81).
- [Banerjee '05] Banerjee, S., & Lavie, A. (2005). METEOR: An automatic metric for MT evaluation with improved correlation with human judgments. In *Proceedings of the acl workshop on intrinsic and extrinsic evaluation measures for machine translation and/or summarization* (pp. 65-72).
- [Papineni '02] Papineni, K., Roukos, S., Ward, T., & Zhu, W. J. (2002, July). BLEU: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics* (pp. 311-318).
- [Wu '16] Wu, Y., Schuster, M., Chen, Z., Le, Q. V., Norouzi, M., Macherey, W., ... & Dean, J. (2016). Google's neural machine translation system: Bridging the gap between human and machine translation. *arXiv preprint arXiv:1609.08144*.
- [Napoles '15] Napoles, C., Sakaguchi, K., Post, M., & Tetreault, J. (2015). Ground truth for grammatical error correction metrics. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)* (pp. 588-593).

- [Sellam '20] Sellam, T., Das, D., & Parikh, A. P. (2020). BLEURT: Learning robust metrics for text generation. arXiv preprint arXiv:2004.04696.
- [Srivastava '22] Srivastava, A., Rastogi, A., Rao, A., Shueb, A. A. M., Abid, A., Fisch, A., ... & Wang, G. (2022). Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. arXiv preprint arXiv:2206.04615.
- [Zellers '19] Zellers, R., Holtzman, A., Bisk, Y., Farhadi, A., & Choi, Y. (2019). Hellaswag: Can a machine really finish your sentence?. arXiv preprint arXiv:1905.07830.
- [Hendrycks '20] Hendrycks, D., Burns, C., Basart, S., Zou, A., Mazeika, M., Song, D., & Steinhardt, J. (2020). Measuring massive multitask language understanding. arXiv preprint arXiv:2009.03300.
- [Simonyan '14] Simonyan, K., & Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. arXiv preprint arXiv:1409.1556.
- [Hubel '59] Hubel, D. H., & Wiesel, T. N. (1959). Receptive fields of single neurones in the cat's striate cortex. *The Journal of physiology*, 148(3), 574-591.
- [Szegedy '15] Szegedy, C., Liu, W., Jia, Y., Sermanet, P., Reed, S., Anguelov, D., ... & Rabinovich, A. (2015). Going deeper with convolutions. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 1-9).
- [Zhang '22] Zhang, A., Lipton, Z. C., Li, M., & Smola, A. J. (2022). Dive into deep learning. <https://d2l.ai>.
- [He '16] He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 770-778).
- [Lofe '15] Lofe, S., & Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *International conference on machine learning* (pp. 448-456). ICML.
- [Ulyanov '17] Ulyanov, D., Vedaldi, A., & Lempitsky, V. (2017). Instance normalization: The missing ingredient for fast stylization. arXiv preprint arXiv:1607.08022.
- [Ba '16] Ba, J. L., Kiros, J. R., & Hinton, G. E. (2016). Layer normalization. arXiv preprint arXiv:1607.06450.
- [Wu '18] Wu, Y., & He, K. (2018). Group normalization. In *Proceedings of the European conference on computer vision (ECCV)* (pp. 3-19).
- [Li '17] Li, H., Xu, Z., Taylor, G., Studer, C., & Goldstein, T. (2017). Visualizing the loss landscape of neural nets. arXiv preprint arXiv:1712.09913.
- [Szegedy '17] Szegedy, C., Ioffe, S., Vanhoucke, V., & Alemi, A. A. (2017). Inception-v4, inception-resnet, and the impact of residual connections on learning. In *The Thirty-first AAAI conference on artificial intelligence*.
- [Dosovitskiy '20] Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., ... & Houlsby, N. (2020). An image is worth 16x16 words: Transformers for image recognition at scale. arXiv preprint arXiv:2010.11929.
- [Girshick '15] Girshick, R., Donahue, J., Darrell, T., & Malik, J. (2015). Rich feature hierarchies for accurate object detection and semantic segmentation. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 580-587).
- [Uijlings '13] Uijlings, J. R., Van De Sande, K. E., Gevers, T., & Smeulders, A. W. (2013). Selective search for object recognition. *International journal of computer vision*, 104(2), 154-171.
- [Girshick '15] Girshick, R. (2015). Fast R-CNN. In *Proceedings of the IEEE international conference on computer vision* (pp. 1440-1448).
- [Ren '16] Ren, S., He, K., Girshick, R., & Sun, J. (2016). Faster R-CNN: Towards real-time object detection with region proposal networks. *Advances in neural information processing systems*, 28, 91-99.
- [Liu '16] Liu, W., Anguelov, D., Erhan, D., Szegedy, C., Reed, S., Fu, C. Y., & Berg, A. C. (2016, October). SSD: Single shot multibox detector. In *European conference on computer vision* (pp. 21-37). Springer, Cham.
- [Redmon '16] Redmon, J., Divvala, S., Girshick, R., & Farhadi, A. (2016). You only look once: Unified, real-time object detection. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 779-788).
- [Redmon '17] Redmon, J., & Farhadi, A. (2017). YOLO9000: better, faster, stronger. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 7263-7271).
- [Redmon '18] Redmon, J., & Farhadi, A. (2018). YOLO v3: An incremental improvement. arXiv preprint arXiv:1804.02767.

- [Lin '17] Lin, T. Y., Dollár, P., Girshick, R., He, K., Hariharan, B., & Belongie, S. (2017). Feature pyramid networks for object detection. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 2117-2125).
- [Ge '21] Ge, Z., Liu, S., Wang, F., Li, Z., & Sun, J. (2021). YoloX: Exceeding yolo series in 2021. *arXiv preprint arXiv:2107.08430*.
- [Bochkovskiy '20] Bochkovskiy, A., Wang, C. Y., & Liao, H. Y. M. (2020). Yolov4: Optimal speed and accuracy of object detection. *arXiv preprint arXiv:2004.10934*.
- [Luo '16] Luo, W., Li, Y., Urtasun, R., & Zemel, R. (2016). Understanding the effective receptive field in deep convolutional neural networks. In *Proceedings of the 30th International Conference on Neural Information Processing Systems* (pp. 4905-4913).
- [Long '15] Long, J., Shelhamer, E., & Darrell, T. (2015). Fully convolutional networks for semantic segmentation. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 3431-3440).
- [He '17] He, K., Gkioxari, G., Dollár, P., & Girshick, R. (2017). Mask R-CNN. In *Proceedings of the IEEE international conference on computer vision* (pp. 2961-2969).
- [Jaderberg '15] Jaderberg, M., Simonyan, K., & Zisserman, A. (2015). Spatial transformer networks. *Advances in neural information processing systems*, 28, 2017-2025.
- [Chen '14] Chen, L. C., Papandreou, G., Kokkinos, I., Murphy, K., & Yuille, A. L. (2014). Semantic image segmentation with deep convolutional nets and fully connected crfs. *arXiv preprint arXiv:1412.7062*.
- [Chen '17 A] Chen, L. C., Papandreou, G., Kokkinos, I., Murphy, K., & Yuille, A. L. (2017). DeepLab: Semantic image segmentation with deep convolutional nets, atrous convolution, and fully connected crfs. *IEEE transactions on pattern analysis and machine intelligence*, 40(4), 834-848.
- [Chen '17 B] Chen, L. C., Papandreou, G., Schroff, F., & Adam, H. (2017). Rethinking Atrous convolution for semantic image segmentation. *arXiv preprint arXiv:1706.05587*.
- [Chen '18] Chen, L. C., Zhu, Y., Papandreou, G., Schroff, F., & Adam, H. (2018). Encoder-decoder with atrous separable convolution for semantic image segmentation. In *Proceedings of the European conference on computer vision (ECCV)* (pp. 801-818).
- [Chollet '17] Chollet, F. (2017). Xception: Deep learning with depthwise separable convolutions. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 1251-1258).
- [Kirillov '23] Kirillov, A., Mintun, E., Ravi, N., Mao, H., Rolland, C., Gustafson, L., ... & Girshick, R. (2023). Segment Anything. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* (pp. 4015-4026).
- [Ravi '24] Ravi, N., Gabeur, V., Hu, Y. T., Hu, R., Ryali, C., Ma, T., ... & Feichtenhofer, C. (2024). SAM 2: Segment Anything in Images and Videos. *arXiv preprint arXiv:2408.00714*.
- [He '22] He, K., Chen, X., Xie, S., Li, Y., Dollár, P., & Girshick, R. (2022). Masked autoencoders are scalable vision learners. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition* (pp. 16000-16009).
- [Mildenhall '21] Mildenhall, B., Srinivasan, P. P., Tancik, M., Barron, J. T., Ramamoorthi, R., & Ng, R. (2021). NeRF: Representing scenes as neural radiance fields for view synthesis. *Communications of the ACM*, 65(1), 99-106.
- [Kerbl '23] Kerbl, B., Kopanas, G., Leimkühler, T., & Drettakis, G. (2023). 3D Gaussian splatting for real-time radiance field rendering. *ACM Transactions on Graphics*, 42(4), 1-14.
- [Snavely '06] Snavely, N., Seitz, S. M., & Szeliski, R. (2006). Photo tourism: exploring photo collections in 3D. In *ACM siggraph 2006 papers* (pp. 835-846).
- [Merrill '10] Merrill, D. G., & Grimshaw, A. S. (2010). Revisiting sorting for GPGPU stream architectures. In *Proceedings of the 19th International Conference on Parallel Architectures and Compilation Techniques* (pp. 545-546).
- [Kling '21] Kling, M. U. (2021). *Quality Land*. Éditions Actes Sud.
- [Masi '18] Masi, I., Wu, Y., Hassner, T., & Natarajan, P. (2018). Deep face recognition: A survey. In *2018 31st IEEE SIBGRAPI conference on graphics, patterns and images (SIBGRAPI)* (pp. 471-478).
- [Taigman '14] Taigman, Y., Yang, M., Ranzato, M. A., & Wolf, L. (2014). Deepface: Closing the gap to human-level performance in face verification. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 1701-1708).

- [Sun '14] Sun, Y., Wang, X., & Tang, X. (2014). Deep learning face representation from predicting 10,000 classes. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 1891-1898).
- [Sun '15] Sun, Y., Wang, X., & Tang, X. (2015). Deeply learned face representations are sparse, selective, and robust. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 2892-2900).
- [Schroff '15] Schroff, F., Kalenichenko, D., & Philbin, J. (2015). Facenet: A unified embedding for face recognition and clustering. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 815-823).
- [Ledig '17] Ledig, C., Theis, L., Huszár, F., Caballero, J., Cunningham, A., Acosta, A., ... & Shi, W. (2017). Photo-realistic single image super-resolution using a generative adversarial network. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 4681-4690.
- [Wang '18] Wang, X., Yu, K., Wu, S., Gu, J., Liu, Y., Dong, C., ... & Change Loy, C. (2018). ESRGAN: Enhanced super-resolution generative adversarial networks. Proceedings of the European Conference on Computer Vision (ECCV), Workshops.
- [Sajjadi '17] Sajjadi, M. S. M., Schölkopf, B., & Hirsch, M. (2017). EnhanceNet: Single image super-resolution through automated texture synthesis. Proceedings of the IEEE International Conference on Computer Vision (ICCV), 4501-4510.
- [Zhang '21] Zhang, Z., Luo, P., Loy, C. C., & Tang, X. (2021). Applications of Human Pose Estimation in Visual Surveillance and Sports: A Survey. Computer Vision and Image Understanding, 206, 103225.
- [Kendall '15] Kendall, A., Grimes, M., & Cipolla, R. (2015). PoseNet: A Convolutional Network for Real-Time 6-DOF Camera Relocalization. *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*.
- [Cao '17] Cao, Z., Simon, T., Wei, S. E., & Sheikh, Y. (2017). Realtime Multi-Person 2D Pose Estimation using Part Affinity Fields. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 7291-7299.
- [Fang '17] Fang, H. S., Xie, S., Tai, Y. W., & Lu, C. (2017). RMPE: Regional Multi-Person Pose Estimation. *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 2334-2343.
- [Alhashim '18] Alhashim, I., & Wonka, P. (2018). High Quality Monocular Depth Estimation via Transfer Learning. *arXiv preprint arXiv:1812.11941*.
- [Laina '16] Laina, I., Rupprecht, C., Belagiannis, V., Tombari, F., & Navab, N. (2016). Deeper Depth Prediction with Fully Convolutional Residual Networks. Proceedings of the 4th International Conference on 3D Vision (3DV), 239-248.
- [Zhu '23] Zhu, Y., Baca, J., Rekabdar, B., & Rawassizadeh, R. (2023). A Survey of AI Music Generation Tools and Models. *arXiv preprint arXiv:2308.12982*.
- [Liu '18] Liu, R., Cornelius, C., Rawassizadeh, R., Peterson, R., & Kotz, D. (2018). Vocal resonance: Using internal body voice for wearable authentication. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, 2(1), 1-23.
- [Arik '18] Arik, S. O., Chrzanowski, M., Coates, A., Duan, G., & Gibson, A. (2018). Deep Voice: Real-time Neural Text-to-Speech. Proceedings of the 34th International Conference on Machine Learning (ICML), 195-204.
- [Oord '16] Oord, A. V. D., Dieleman, S., Zen, H., Simonyan, K., Vinyals, O., Graves, A., ... & Kavukcuoglu, K. (2016). Wavenet: A generative model for raw audio. *arXiv preprint arXiv:1609.03499*.
- [Oord '18] Oord, A., Li, Y., Babuschkin, I., Simonyan, K., Vinyals, O., Kavukcuoglu, K., ... & Hassabis, D. (2018). Parallel wavenet: Fast high-fidelity speech synthesis. In International conference on machine learning (pp. 3918-3926).
- [Wang '17] Wang, Y., Skerry-Ryan, R. J., Stanton, D., Wu, Y., Weiss, R. J., Jaitly, N., ... & Saurous, R. A. (2017). Tacotron: Towards end-to-end speech synthesis. *arXiv preprint arXiv:1703.10135*.
- [Srivastava '15] Srivastava, R. K., Greff, K., & Schmidhuber, J. (2015). Training very deep networks. *Advances in neural information processing systems*, 28.
- [Shen '18] Shen, J., Pang, R., Weiss, R. J., Schuster, M., Jaitly, N., Yang, Z., ... & Wu, Y. (2018). Natural TTS synthesis by conditioning wavenet on mel spectrogram predictions. In 2018 IEEE international conference on acoustics, speech and signal processing (ICASSP) (pp. 4779-4783).
- [Chorowski '15] Chorowski, J. K., Bahdanau, D., Serdyuk, D., Cho, K., & Bengio, Y. (2015). Attention-based models for speech recognition. *Advances in Neural Information Processing Systems*, 28.

[Jang '16] Jang, E., Gu, S., & Poole, B. (2016). Categorical reparameterization with gumbel-softmax. arXiv preprint arXiv:1611.01144.

[Schneider '19] Schneider, S., Baevski, A., Collobert, R., & Auli, M. (2019). wav2vec: Unsupervised pre-training for speech recognition. arXiv preprint arXiv:1904.05862.

[Baevski '20] Baevski, A., Zhou, Y., Mohamed, A., & Auli, M. (2020). wav2vec 2.0: A framework for self-supervised learning of speech representations. *Advances in neural information processing systems*, 33, 12449-12460.

[Radford '23] Radford, A., Kim, J. W., Xu, T., Brockman, G., McLeavey, C., & Sutskever, I. (2023). Robust speech recognition via large-scale weak supervision. *In International Conference on Machine Learning* (pp. 28492-28518).

———— Chapter 13 ————

[Skinner '65] Skinner, B. F. (1965). Science and human behavior (No. 92904). Simon and Schuster.

[Sutton '20] Sutton, R. S., & Barto, A. G. (2020). Reinforcement learning: An introduction. MIT press.

[Silver '17] Silver, D., Schrittwieser, J., Simonyan, K., Antonoglou, I., Huang, A., Guez, A., ... & Hassabis, D. (2017). Mastering the game of go without human knowledge. *nature*, 550(7676), 354-359.

[Silver '18] Silver, D., Hubert, T., Schrittwieser, J., Antonoglou, I., Lai, M., Guez, A., ... & Hassabis, D. (2018). A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play. *Science*, 362(6419), 1140-1144.

[Ouyang '22] Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C. L., Mishkin, P., ... & Lowe, R. (2022). Training language models to follow instructions with human feedback. *arXiv preprint arXiv:2203.02155*.

[Thompson '33] Thompson, W. R. (1933). On the likelihood that one unknown probability exceeds another in view of the evidence of two samples. *Biometrika*, 25(3-4), 285-294.

[Winder '20] Winder, P. (2020). Reinforcement learning. O'Reilly Media.

[Sutton '98] Richard S. (1988). "Learning to predict by the methods of temporal differences". *Machine Learning*. 3 (1): 9–44.

[Watkins '89] Watkins, C.J.C.H. (1989). Learning from Delayed Rewards. Ph.D. thesis, University of Cambridge.

[Sutton '91] Sutton, R. S. (1991). Dyna, an integrated architecture for learning, planning, and reacting. *ACM Sigart Bulletin*, 2(4), 160-163.

[Peng '94] Peng, J., & Williams, R. J. (1994). Incremental multi-step Q-learning. In *Machine Learning Proceedings 1994* (pp. 226-232). Morgan Kaufmann.

[Morales '20] Morales, M. (2020). Grokking deep reinforcement learning. Manning Publications.

[Mnih '13] Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., & Riedmiller, M. (2013). Playing atari with deep reinforcement learning. arXiv preprint arXiv:1312.5602.

[Mnih '15] Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., ... & Hassabis, D. (2015). Human-level control through deep reinforcement learning. *nature*, 518(7540), 529-533.

[Riedmiller '05] Riedmiller, M. (2005). Neural fitted Q iteration—first experiences with a data efficient neural reinforcement learning method. In *European conference on machine learning* (pp. 317-328). Springer, Berlin, Heidelberg.

[Lin '92] Lin, L. J. (1992). Self-improving reactive agents based on reinforcement learning, planning and teaching. *Machine learning*, 8(3), 293-321.

[Van Hasselt '16] Van Hasselt, H., Guez, A., & Silver, D. (2016). Deep reinforcement learning with double q-learning. In *Proceedings of the AAAI conference on artificial intelligence* (Vol. 30, No. 1).

[Van Hasselt '10] Van Hasselt, H. (2010). Double Q-learning. *Advances in Neural Information Processing Systems*, 23.

[Wang '16] Wang, Z., Schaul, T., Hessel, M., Hasselt, H., Lanctot, M., & Freitas, N. (2016). Dueling network architectures for deep reinforcement learning. In *International conference on machine learning* (pp. 1995-2003). PMLR.

[Schaul '16] Schaul, T., Quan, J., Antonoglou, I., & Silver, D. (2015). Prioritized experience replay. arXiv preprint arXiv:1511.05952.

[Williams '92] Williams, R. J. (1992). Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine learning*, 8(3), 229-256.

- [Schulman '15] Schulman, J., Levine, S., Abbeel, P., Jordan, M., & Moritz, P. (2015, June). Trust region policy optimization. In International conference on machine learning (pp. 1889-1897). PMLR.
- [Amari '98] Amari, S. I. (1998). Natural gradient works efficiently in learning. *Neural computation*, 10(2), 251-276.
- [Stiefel '52] Stiefel, E. (1952). Methods of conjugate gradients for solving linear systems. *Journal of Research of the National Bureau of Standards*, 49, 409-435.
- [Chen '16] Chen, T., Xu, B., Zhang, C., & Guestrin, C. (2016). Training deep nets with Sublinear Memory cost. arXiv preprint arXiv:1604.0617
- [Schulman '17] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms. arXiv preprint arXiv:1707.06347.
- [Mnih '16] Mnih, V., Badia, A. P., Mirza, M., Graves, A., Lillicrap, T., Harley, T., ... & Kavukcuoglu, K. (2016). Asynchronous Methods for Deep Reinforcement Learning. In International Conference on Machine Learning (pp. 1928-1937). PMLR.
- [Lillicrap '15] Lillicrap, T. P., Hunt, J. J., Pritzel, A., Heess, N., Erez, T., Tassa, Y., ... & Wierstra, D. (2015). Continuous control with deep reinforcement learning. arXiv preprint arXiv:1509.02971.
- [Uhlenbeck '30] Uhlenbeck, G. E and Ornstein, L.S. (1930) On the theory of the brownian motion. *Physical review*, 36(5):823.
- [Fujimoto '18] Fujimoto, S., Hoof, H., & Meger, D. (2018). Addressing function approximation error in actor-critic methods. In International Conference on Machine Learning (pp. 1587-1596). PMLR.
- [Haarnoja '18] Haarnoja, T., Zhou, A., Abbeel, P., & Levine, S. (2018). Soft Actor-Critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. In International conference on machine learning (pp. 1861-1870). PMLR.
- [Bengio '13] Bengio, Y., Léonard, N., & Courville, A. (2013). Estimating or propagating gradients through stochastic neurons for conditional computation. arXiv preprint arXiv:1308.3432.
- [Team '24] Team, O. M., Ghosh, D., Walke, H., Pertsch, K., Black, K., Mees, O., ... & Levine, S. (2024). Octo: An Open-Source Generalist Robot Policy. arXiv preprint arXiv:2405.12213.
- [Hafner '19] Hafner, D., Lillicrap, T., Ba, J., & Norouzi, M. (2019). Dream to control: Learning behaviors by latent imagination. arXiv preprint arXiv:1912.01603.
- [Hafner '20] Hafner, D., Lillicrap, T., Norouzi, M., & Ba, J. (2020). Mastering atari with discrete world models. arXiv preprint arXiv:2010.02193.
- [Hafner '23] Hafner, D., Pasukonis, J., Ba, J., & Lillicrap, T. (2023). Mastering Diverse Domains through World Models. arXiv preprint arXiv:2301.04104.
- [Ziyin '20] Ziyin, L., Wang, Z. T., & Ueda, M. (2020). LaProp: Separating momentum and adaptivity in adam. arXiv preprint arXiv:2002.04839.

———— Chapter 14 ————

- [Sedgewick '11] Sedgewick, R., & Wayne, K. (2011). *Algorithms, 4th Edition*. Addison-Wesley Professional.
- [Cormen '22] Cormen, T. H., Leiserson, C. E., Rivest, R. L., & Stein, C. (2009). *Introduction to algorithms, 4th Edition*. MIT press.
- [Ziv '78] Ziv, J., & Lempel, A. (1978). Compression of Individual Sequences via Variable-Rate Coding. *IEEE transactions on Information Theory*, 24(5), 530-536.
- [Welch '84] Welch, T. A. (1984). Technique for high-performance data compression. *Computer*, (52).
- [Rawassizadeh '23] Rawassizadeh, R., & Rong, Y. (2023). ODSearch: Fast and Resource Efficient On-device Natural Language Search for Fitness Trackers' Data. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, 6(4), 1-25.
- [Malekijoo '21] Malekijoo, A., Fadaeieslam, M. J., Malekijoo, H., Homayounfar, M., Alizadeh-Shabdiz, F., & Rawassizadeh, R. (2021). Fedzip: A compression framework for communication-efficient federated learning. arXiv preprint arXiv:2102.01593.
- [Gage '94] Gage, P. (1994). A new algorithm for data compression. *The C Users Journal*, 12(2), 23-38.
- [Huffman '52] Huffman, D. A. (1952). A method for the construction of minimum-redundancy codes. *Proceedings of the IRE*, 40(9), 1098-1101.

- [Linde '80] Linde, Y., Buzo, A., & Gray, R. (1980). An algorithm for vector quantizer design. *IEEE Transactions on communications*, 28(1), 84-95.
- [Horowitz '14] Horowitz, M. (2014). 1.1 Computing's Energy Problem (and what we can do about it). In 2014 IEEE International Solid-State Circuits Conference Digest of Technical Papers (ISSCC) (pp. 10-14). IEEE.
- [Krishnamoorthi '18] Krishnamoorthi, R. (2018). Quantizing deep convolutional networks for efficient inference: A whitepaper. *arXiv preprint arXiv:1806.08342*.
- [Jacob '18] Jacob, B., Kligys, S., Chen, B., Zhu, M., Tang, M., Howard, A., ... & Kalenichenko, D. (2018). Quantization and training of neural networks for efficient integer-arithmetic-only inference. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 2704-2713).
- [Courbariaux '15] Courbariaux, M., Bengio, Y., & David, J. P. (2015). Binaryconnect: Training deep neural networks with binary weights during propagations. *Advances in neural information processing systems*, 28.
- [Li '16] Li, F., Zhang, B., & Liu, B. (2016). Ternary weight networks. *arXiv preprint arXiv:1605.04711*.
- [Wang '23] Wang, H., Ma, S., Dong, L., Huang, S., Wang, H., Ma, L., ... & Wei, F. (2023). Bitnet: Scaling 1-bit transformers for large language models. *arXiv preprint arXiv:2310.11453*.
- [Alizadeh '20] Alizadeh, M., Behboodi, A., van Baalen, M., Louizos, C., Blankevoort, T., & Welling, M. (2020). Gradient ℓ_1 regularization for quantization robustness. *arXiv preprint arXiv:2002.07520*.
- [LeCun '89] LeCun, Yann, John Denker, and Sara Solla. "Optimal brain damage." *Advances in neural information processing systems* 2 (1989).
- [Hassibi '93] Hassibi, B., Stork, D. G., & Wolff, G. J. (1993). Optimal brain surgeon and general network pruning. In *IEEE international conference on neural networks* (pp. 293-299).
- [Reed '93] Reed, R. (1993). Pruning algorithms-a survey. *IEEE transactions on Neural Networks*, 4(5), 740-747.
- [Hoefler '21] Hoefler, T., Alistarh, D., Ben-Nun, T., Dryden, N., & Peste, A. (2021). Sparsity in deep learning: Pruning and growth for efficient inference and training in neural networks. *The Journal of Machine Learning Research*, 22(1), 10882-11005.
- [Li '16] Li, H., Kadav, A., Durdanovic, I., Samet, H., & Graf, H. P. (2016). Pruning filters for efficient convnets. *arXiv preprint arXiv:1608.08710*.
- [Molchanov '16] Molchanov, P., Tyree, S., Karras, T., Aila, T., & Kautz, J. (2016). Pruning convolutional neural networks for resource efficient inference. *arXiv preprint arXiv:1611.06440*.
- [Lin '17] Lin, J., Rao, Y., Lu, J., & Zhou, J. (2017). Runtime neural pruning. *Advances in neural information processing systems*, 30.
- [Williams '95] Williams, P. M. (1995). Bayesian regularization and pruning using a Laplace prior. *Neural computation*, 7(1), 117-143.
- [Wu '21] Wu, D., & Wang, Y. (2021). Adversarial neuron pruning purifies backdoored deep models. *Advances in Neural Information Processing Systems*, 34, 16913-16925.
- [Frankle '18] Frankle, J., & Carbin, M. (2018). The lottery ticket hypothesis: Finding sparse, trainable neural networks. *arXiv preprint arXiv:1803.03635*.
- [Lu '24] Lu, H., Alemi, M., & Rawassizadeh, R. (2024). The Impact of Quantization and Pruning on Deep Reinforcement Learning Models. *arXiv preprint arXiv:2407.04803*.
- [Hu '21] Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., ... & Chen, W. (2021). Lora: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*.
- [Jiang '23] Jiang, A. Q., Sablayrolles, A., Mensch, A., Bamford, C., Chaplot, D. S., Casas, D. D. L., ... & Sayed, W. E. (2023). Mistral 7B. *arXiv preprint arXiv:2310.06825*.
- [Dao '22] Dao, T., Fu, D., Ermon, S., Rudra, A., & Ré, C. (2022). FlashAttention: Fast and memory-efficient exact attention with io-awareness. *Advances in Neural Information Processing Systems*, 35, 16344-16359.
- [Dao '23] Dao, T. (2023). FlashAttention-2: Faster attention with better parallelism and work partitioning. *arXiv preprint arXiv:2307.08691*.
- [Shazeer '19] Shazeer, N. (2019). Fast transformer decoding: One write-head is all you need. *arXiv preprint arXiv:1911.02150*.
- [Chowdhery '23] Chowdhery, A., Narang, S., Devlin, J., Bosma, M., Mishra, G., Roberts, A., ... & Fiedel, N. (2023). Palm: Scaling language modeling with pathways. *Journal of Machine Learning Research*, 24(240), 1-113.

- [Ainslie '23] Ainslie, J., Lee-Thorp, J., de Jong, M., Zemlyanskiy, Y., Lebrón, F., & Sanghai, S. (2023). GQA: Training generalized multi-query transformer models from multi-head checkpoints. arXiv preprint arXiv:2305.13245.
- [Beltagy '20] Beltagy, I., Peters, M. E., & Cohan, A. (2020). Longformer: The long-document transformer. arXiv preprint arXiv:2004.05150.
- [Shui '19] Shui, C., Abbasi, M., Robitaille, L. É., Wang, B., & Gagné, C. (2019). A principled approach for learning task similarity in multitask learning. arXiv preprint arXiv:1903.09109.
- [Aoki '22] Aoki, R., Tung, F., & Oliveira, G. L. (2022). Heterogeneous multi-task learning with expert diversity. *IEEE/ACM Transactions on Computational Biology and Bioinformatics*, 19(6), 3093-3102.
- [Schmidhuber '87] Schmidhuber, J. (1987). Evolutionary principles in self-referential learning, or on learning how to learn: the meta-meta-... hook (Doctoral dissertation, Technische Universität München).
- [Santoro '16] Santoro, A., Bartunov, S., Botvinick, M., Wierstra, D., & Lillicrap, T. (2016). Meta-learning with memory-augmented neural networks. In *International Conference on Machine Learning* (pp. 1842-1850). PMLR.
- [Finn '17] Finn, C., Abbeel, P., & Levine, S. (2017). Model-Agnostic Meta-Learning for fast adaptation of deep networks. In *International conference on machine learning* (pp. 1126-1135). PMLR.
- [Weinshall '18] Weinshall, D., Cohen, G., & Amir, D. (2018). Curriculum learning by transfer learning: Theory and experiments with deep networks. In *International Conference on Machine Learning* (pp. 5238-5246). PMLR.
- [Karimi '22] Karimi, F., Mehrpanah, A., & Rawassizadeh, R. (2022). LightDepth: A Resource Efficient Depth Estimation Approach for Dealing with Ground Truth Sparsity via Curriculum Learning. arXiv preprint arXiv:2211.08608.
- [Konečný '18] Konečný, J., McMahan, H. B., Ramage, D., & Richtárik, P. (2016). Federated optimization: Distributed machine learning for on-device intelligence. arXiv preprint arXiv:1610.02527.
- [McMahan '17] McMahan, B., Moore, E., Ramage, D., Hampson, S., & y Arcas, B. A. (2017). Communication-efficient learning of deep networks from decentralized data. In *Artificial intelligence and statistics* (pp. 1273-1282). PMLR.
- [Acar '21] Acar, D. A. E., Zhao, Y., Navarro, R. M., Mattina, M., Whatmough, P. N., & Saligrama, V. (2021). Federated learning based on dynamic regularization. arXiv preprint arXiv:2111.04263.
- [Rawassizadeh '18] Rawassizadeh, R., Pierson, T. J., Peterson, R., & Kotz, D. (2018). NoCloud: Exploring network disconnection through on-device data analysis. *IEEE Pervasive Computing*, 17(1), 64-74.
- [Cheng '17] Cheng, Y., Wang, D., Zhou, P., & Zhang, T. (2017). A survey of model compression and acceleration for deep neural networks. arXiv preprint arXiv:1710.09282.
- [Liu '22] Liu, Y., Cao, J., Li, B., Hu, W., Ding, J., & Li, L. (2022). Cross-Architecture Knowledge Distillation. In *Proceedings of the Asian Conference on Computer Vision* (pp. 3396-3411).
- [Wang '18] Wang, H., Zhao, H., Li, X., & Tan, X. (2018). Progressive Blockwise Knowledge Distillation for Neural Network Acceleration. In *IJCAI* (pp. 2769-2775).
- [Li '20] Li, C., Peng, J., Yuan, L., Wang, G., Liang, X., Lin, L., & Chang, X. (2020). Block-wisely supervised neural architecture search with knowledge distillation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 1989-1998).
- [Romero '14] Romero, A., Ballas, N., Kahou, S. E., Chassang, A., Gatta, C., & Bengio, Y. (2014). Fitnets: Hints for thin deep nets. arXiv preprint arXiv:1412.6550.
- [Chen '17] Chen, G., Choi, W., Yu, X., Han, T., & Chandraker, M. (2017). Learning efficient object detection models with knowledge distillation. *Advances in neural information processing systems*, 30.
- [Elsken '19] Elsken, T., Metzen, J. H., & Hutter, F. (2019). Neural architecture search: A survey. *The Journal of Machine Learning Research*, 20(1), 1997-2017.
- [Gholami '21] Gholami, A., Kim, S., Dong, Z., Yao, Z., Mahoney, M. W., & Keutzer, K. (2021). A survey of quantization methods for efficient neural network inference. arXiv preprint arXiv:2103.13630.
- [Hoefler '21] Hoefler, T., Alistarh, D., Ben-Nun, T., Dryden, N., & Peste, A. (2021). Sparsity in deep learning: Pruning and growth for efficient inference and training in neural networks. *The Journal of Machine Learning Research*, 22(1), 10882-11005.

- [Howard '17] Howard, A. G., Zhu, M., Chen, B., Kalenichenko, D., Wang, W., Weyand, T., ... & Adam, H. (2017). Mobilenets: Efficient convolutional neural networks for mobile vision applications. *arXiv preprint arXiv:1704.04861*.
- [Iandola '16] Iandola, F. N., Han, S., Moskewicz, M. W., Ashraf, K., Dally, W. J., & Keutzer, K. (2016). SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and < 0.5 MB model size. *arXiv preprint arXiv:1602.07360*.
- [Tan '19] Tan, M., & Le, Q. (2019). Efficientnet: Rethinking model scaling for convolutional neural networks. In *International conference on machine learning* (pp. 6105-6114). PMLR.
- [Mehta '21] Mehta, S., & Rastegari, M. (2021). Mobilevit: light-weight, general-purpose, and mobile-friendly vision transformer. *arXiv preprint arXiv:2110.02178*.

———— Chapter 15 ————

- [Salathe '11] Salathé, M., & Khandelwal, S. (2011). Assessing vaccination sentiments with online social media: implications for infectious disease dynamics and control. *PLoS computational biology*, 7(10), e1002199.
- [Collier '11] Collier, N., Son, N. T., & Nguyen, N. M. (2011). OMG U got flu? Analysis of shared health messages for bio-surveillance. *Journal of biomedical semantics*, 2(5), 1-10.
- [Chen '11] Chen, Y., Fay, S., & Wang, Q. (2011). The role of marketing in social media: How online consumer reviews evolve. *Journal of interactive marketing*, 25(2), 85-94.
- [Gan '11] Gan, Y., Wang, T., Javaheri, A., Momeni-Ortner, E., Dehghani, M., Hosseinzadeh, M., & Rawassizadeh, R. (2021). 11 Years with Wearables: Quantitative Analysis of Social Media, Academia, News Agencies, and Lead User Community from 2009-2020 on Wearable Technologies. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, 5(1), 1-26.
- [Bradshaw '17] Bradshaw, S., & Howard, P. (2017). Troops, trolls and troublemakers: A global inventory of organized social media manipulation. *Computational Propaganda Research Project*.
- [Jarník '30] Jarník, V. (1930), "O jistém problému minimálním" [About a certain minimal problem], *Práce Moravské Přírodovědecké Společnosti* (in Czech), 6 (4): 57–63.
- [Prim '57] Prim, R. C. (1957). Shortest connection networks and some generalizations. *The Bell System Technical Journal*, 36(6), 1389-1401.
- [Kruskal '56] Kruskal, J. B. (1956). "On the shortest spanning subtree of a graph and the traveling salesman problem". *Proceedings of the American Mathematical Society*. 7 (1): 48–50.
- [Dijkstra '59] Dijkstra, E. W. (1959). "A note on two problems in connexion with graphs" (PDF). *Numerische Mathematik*. 1: 269–271.
- [Hart '68] Hart, P. E., Nilsson, N. J., & Raphael, B. (1968). A formal basis for the heuristic determination of minimum cost paths. *IEEE Transactions on Systems Science and Cybernetics*, 4(2), 100-107.
- [Kuhn '55] Kuhn, H. W. (1955). The Hungarian method for the assignment problem. *Naval research logistics quarterly*, 2(1-2), 83-97.
- [Hopcroft '73] Hopcroft, J. E., & Karp, R. M. (1973). An $n^{5/2}$ algorithm for maximum matchings in bipartite graphs. *SIAM Journal on Computing*, 2(4), 225-231.
- [Bellman '58] Bellman, Richard (1958). "On a routing problem". *Quarterly of Applied Mathematics*. 16: 87–90.
- [Ford '56] Ford, Lester R. Jr. (August 14, 1956). *Network Flow Theory*. Paper P-923. Santa Monica, California: RAND Corporation.
- [Morselli '08] Morselli, C., & Roy, J. (2008). Brokerage qualifications in ringing operations. *Criminology*, 46(1), 71-98.
- [Kazerani '09] Kazerani, A., & Winter, S. (2009). Can betweenness centrality explain traffic flow. In *12th AGILE international conference on geographic information science* (pp. 1-9).
- [Bonacich '87] Bonacich, P. (1987). Power and centrality: A family of measures. *American journal of sociology*, 92(5), 1170-1182.
- [Page '99] Page, L., Brin, S., Motwani, R., & Winograd, T. (1999). *The PageRank citation ranking: Bringing order to the web*. Stanford InfoLab.
- [Kleinberg '99] Kleinberg, J. M. (1999). Authoritative sources in a hyperlinked environment. *Journal of the ACM (JACM)*, 46(5), 604-632.

- [Cheeger '70] Cheeger, J. (1970). A lower bound for the smallest eigenvalue of the Laplacian, Problems in analysis (Papers dedicated to Salomon Bochner, 1969).
- [Donath '72] Donath, W. E., & Hoffman, A. J. (1972). Algorithms for partitioning of graphs and computer logic based on eigenvectors of connection matrices. *IBM Technical Disclosure Bulletin*, 15(3), 938-944.
- [Fiedler '73] Fiedler, M. (1973). Algebraic connectivity of graphs. *Czechoslovak Mathematical Journal*, 23(2), 298-305.
- [Blondel '08] Blondel, V. D., Guillaume, J. L., Lambiotte, R., & Lefebvre, E. (2008). Fast unfolding of communities in large networks. *Journal of statistical mechanics: theory and experiment*, 2008(10), P10008.
- [Traag '19] Traag, V. A., Waltman, L., & van Eck, N. J. (2019). From Louvain to Leiden: guaranteeing well-connected communities. *Scientific reports*, 9(1), 1-12.
- [Traag '11] Traag, V. A., Van Dooren, P., & Nesterov, Y. (2011). Narrow scope for resolution-limit-free community detection. *Physical Review E*, 84(1), 016114.
- [Bruna '14] Bruna, W. Zaremba, and Y. Szlam, A. and LeCun (2014). Spectral networks and locally connected networks on graphs. In International Conference on Learning Representations (ICLR).
- [Merkwirth '05] Merkwirth, C., & Lengauer, T. (2005). Automatic generation of complementary descriptors with molecular graph networks. *Journal of chemical information and modeling*, 45(5), 1159-1168.
- [Hamilton '20] Hamilton, W. L. (2020). Graph representation learning. *Synthesis Lectures on Artificial Intelligence and Machine Learning*, 14(3), 1-159.
- [Grattarola '22] Grattarola, D., Zambon, D., Bianchi, F. M., & Alippi, C. (2022). Understanding pooling in graph neural networks. *IEEE Transactions on Neural Networks and Learning Systems*.
- [Lee '19] Lee, J., Lee, I., & Kang, J. (2019). Self-attention graph pooling. In International conference on machine learning (pp. 3734-3743). PMLR.
- [Liu '22] Liu, C., Zhan, Y., Li, C., Du, B., Wu, J., Hu, W., ... & Tao, D. (2022). Graph pooling for graph neural networks: Progress, challenges, and opportunities. *arXiv preprint arXiv:2204.07321*.
- [Kipf '16] Kipf, T. N., & Welling, M. (2016). Semi-supervised classification with graph convolutional networks. *arXiv preprint arXiv:1609.02907*.
- [Veličković '17] Veličković, P., Cucurull, G., Casanova, A., Romero, A., Lio, P., & Bengio, Y. (2017). Graph attention networks. *arXiv preprint arXiv:1710.10903*.
- [Hamilton '17] Hamilton, W., Ying, Z., & Leskovec, J. (2017). Inductive representation learning on large graphs. *Advances in neural information processing systems*, 30.
- [Zhou '20] Zhou, J., Cui, G., Hu, S., Zhang, Z., Yang, C., Liu, Z., ... & Sun, M. (2020). Graph neural networks: A review of methods and applications. *AI open*, 1, 57-81.
- [Malkov '18] Malkov, Y. A., & Yashunin, D. A. (2018). Efficient and robust approximate nearest neighbor search using hierarchical navigable small-world graphs. *IEEE transactions on pattern analysis and machine intelligence*, 42(4), 824-836.

———— Chapter 16 ————

- [Griffin '84] Griffin, D., & Lim, J. (1984). Signal estimation from modified short-time Fourier transform. *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 32(2), 236-243.
- [Metropolis '49] Metropolis, N., & Ulam, S. (1949). The Monte Carlo method. *Journal of the American Statistical Association*. 44 (247): 335–341.
- [Metropolis '53] Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., & Teller, E. (1953). Equation of state calculations by fast computing machines. *The journal of Chemical Physics*, 21(6), 1087-1092.
- [Hastings '70] Hastings, W. K. (1970). Monte Carlo sampling methods using Markov chains and their applications.
- [Geman '84] Geman, S.; Geman, D. (1984). "Stochastic Relaxation, Gibbs Distributions, and the Bayesian Restoration of Images". *IEEE Transactions on Pattern Analysis and Machine Intelligence*. 6 (6): 721–741.
- [Kloek '84] Kloek, T., & Van Dijk, H. K. (1978). Bayesian estimates of equation system parameters: an application of integration by Monte Carlo. *Econometrica: Journal of the Econometric Society*, 1-19.
- [Wan '22] Wan, Z., Zhang, B., Chen, D., Zhang, P., Wen, F., & Liao, J. (2022). Old Photo Restoration via Deep Latent Space Translation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 45(2), 2071-2087.

- [Perlin '85] Perlin, K. (1985). An image synthesizer. *ACM Siggraph Computer Graphics*, 19(3), 287-296.
- [Perlin '01] Perlin, K. (2001). *Noise hardware*. In *Real-Time Shading SIGGRAPH Course Notes*, 6, 26.
- [Worley '96] Worley, S. (1996). A cellular texture basis function. In *Proceedings of the 23rd Annual Conference on Computer Graphics and Interactive Techniques* (pp. 291-294).
- [Butterworth '30] Butterworth, S. (1930). On the theory of filter amplifiers. *Wireless Engineer*, 7(6), 536-541.
- [Wiener '49] Wiener, N. (1949). Extrapolation, interpolation, and smoothing of stationary time series: with engineering applications. *The MIT Press*.
- [Kalman '60] Kalman, R. E. (1960). A New Approach to Linear Filtering and Prediction Problems. *Transactions of the ASME--Journal of Basic Engineering*, 82(D), 35-45.
- [Lin '17] Lin, T. Y., Goyal, P., Girshick, R., He, K., & Dollár, P. (2017). Focal loss for dense object detection. In *Proceedings of the IEEE International Conference on Computer Vision* (pp. 2980-2988).
- [Chawla '02] Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16, 321-357.
- [Bunkhumpornpat '09] Bunkhumpornpat, C., Sinapiromsaran, K., & Lursinsap, C. (2009). Safe-level-smote: Safe-level-synthetic minority over-sampling technique for handling the class imbalanced problem. In *Advances in Knowledge Discovery and Data Mining: 13th Pacific-Asia conference, PAKDD 2009*.
- [He '08] He, H., Bai, Y., Garcia, E. A., & Li, S. (2008). ADASYN: Adaptive synthetic sampling approach for imbalanced learning. In *2008 IEEE International Joint Conference on Neural Networks* (pp. 1322-1328).
- [Bissoto '21] Bissoto, A., Valle, E., & Avila, S. (2021). Gan-based data augmentation and anonymization for skin-lesion analysis: A critical review. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition* (pp. 1847-1856).
- [Yu '23] Yu, X., Li, G., Lou, W., Liu, S., Wan, X., Chen, Y., & Li, H. (2023). Diffusion-based data augmentation for nuclei image segmentation. In *International Conference on Medical Image Computing and Computer-Assisted Intervention* (pp. 592-602).
- [Wan '21] Wan, Z., Zhang, Y., & He, H. (2017). Variational autoencoder based synthetic data generation for imbalanced learning. In *2017 IEEE Symposium Series on Computational Intelligence (SSCI)* (pp. 1-7).
- [Batista '23] Batista, G. E., & Monard, M. C. (2003). An analysis of four missing data treatment methods for supervised learning. *Applied Artificial Intelligence*, 17(5-6), 519-533.
- [Rawassizadeh '19] Rawassizadeh, R., Keshavarz, H., & Pazzani, M. (2019). Ghost imputation: Accurately reconstructing missing data of the off period. *IEEE Transactions on Knowledge and Data Engineering*, 32(11), 2185-2197.
- [Stekhoven '12] Stekhoven, D. J., & Bühlmann, P. (2012). MissForest - non-parametric missing value imputation for mixed-type data. *Bioinformatics*, 28(1), 112-118.
- [Oliver '90] Oliver, M. A., & Webster, R. (1990). Kriging: a method of interpolation for geographical information systems. *International Journal of Geographical Information System*, 4(3), 313-332.
- [Hardy '71] Hardy, R. L. (1971). Multiquadric equations of topography and other irregular surfaces. *Journal of Geophysical Research*, 76(8), 1905-1915.
- [Liu '08] Liu, F. T., Ting, K. M., & Zhou, Z. H. (2008). Isolation forest. In *8th IEEE International Conference on Data Mining* (pp. 413-422).
- [Schölkopf '01] Schölkopf, B., Platt, J. C., Shawe-Taylor, J., Smola, A. J., & Williamson, R. C. (2001). Estimating the Support of a High-Dimensional Distribution. *Neural Computation*, 13(7), 1443-1471.
- [Breunig '00] Breunig, M. M., Kriegel, H. P., Ng, R. T., & Sander, J. (2000). LOF: identifying density-based local outliers. In *Proceedings of the 2000 ACM SIGMOD international conference on Management of data* (pp. 93-104).
- [Fischler '81] Fischler, M. A., & Bolles, R. C. (1981). Random Sample Consensus: A Paradigm for Model Fitting with Applications to Image Analysis and Automated Cartography. *Communications of the ACM*, 24(6), 381-395.
- [Bifet '07] Bifet, A., & Gavaldà, R. (2007). Learning from time-changing data with adaptive windowing. In *Proceedings of the 2007 SIAM International Conference on Data Mining* (pp. 443-448).
- [Cohen '60] Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and psychological measurement*, 20(1), 37-46.

- [Fleiss '71] Fleiss, J. L. (1971). Measuring nominal scale agreement among many raters. *Psychological bulletin*, 76(5), 378.
- [Krippendorff '18] Krippendorff, K. (2018). Content analysis: An introduction to its methodology. Sage publications.
- [Lewis '20] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., ... & Kiela, D. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. *Advances in Neural Information Processing Systems*, 33, 9459-9474.
- [Ji '23] Ji, X., Sungu-Eryilmaz, Y., Momeni, E., & Rawassizadeh, R. (2023). Speeding Up Question Answering Task of Language Models via Inverted Index. *The Generative AI for Pervasive Computing Symposium (GenAI4PC)*, in conjunction with ACM Interactive, Mobile, Wearable and Ubiquitous Technologies. arXiv:2210.13578.
- [Wei '22] Wei, J., Wang, X., Schuurmans, D., Bosma, M., Xia, F., Chi, E., ... & Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35, 24824-24837.
- [Brown '20] Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., ... & Amodei, D. (2020). Language Models are Few-shot Learners. *Advances in Neural Information Processing Systems*, 33, 1877-1901.
- [Wang '22] Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., ... & Zhou, D. (2022). Self-consistency Improves the Chain of Thought Reasoning in Language Models. *arXiv preprint arXiv:2203.11171*.
- [Yao '23] Yao, S., Yu, D., Zhao, J., Shafran, I., Griffiths, T. L., Cao, Y., & Narasimhan, K. (2023). Tree of Thoughts: Deliberate Problem Solving with Large Language Models. *arXiv preprint arXiv:2305.10601*.
- [Long '23] Long, J. (2023). Large Language Model Guided Tree-of-Thought. *arXiv preprint arXiv:2305.08291*.
- [Kahneman '11] Kahneman, D. (2011). Thinking, fast and slow. MacMillan.
- [Bansal '18] Bansal, N., Chen, X., & Wang, Z. (2018). Can we gain more from orthogonality regularizations in training deep networks?. *Advances in Neural Information Processing Systems*, 31.
- [Lundberg '17] Lundberg, S. M., & Lee, S. I. (2017). A unified approach to interpreting model predictions. *Advances in neural information processing systems*, 30.
- [Shapley '51] Shapley, L. S. (1953). A value for n-person games. In *H. W. Kuhn & A. W. Tucker (Eds.), Contributions to the Theory of Games* (Vol. 2, pp. 307-317). Princeton University Press.
- [Ribeiro '16] Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). "Why should I trust you?" Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD international conference on Knowledge Discovery and Data mining* (pp. 1135-1144).