

Table of Content

Part i: Preliminary Concepts

Chapter 1: Introduction and Basic Concepts	3
Ultra Short History of AI and Machine Learning	3
What is Data Science, Machine Learning, Data Mining, and Artificial Intelligence?	4
Why is data science important?	5
What is an algorithm?	6
Evaluating a Machine Learning Algorithm	7
Computational Complexity and Efficiency	7
Algorithms Runtime	9
Accuracy	10
The Balance between Accuracy and Efficiency	12
Types of Input Dataset	13
Tabular Dataset	13
Temporal Datasets	14
Data Streams	15
Graph Datasets	15
Text	17
Image and Video	17
Audio	18
Numerical Data Types	18
Tasks that Machine Learning Can Perform	21
Unsupervised Learning	22
Clustering	22
Association Rule and Sequence Mining	22
Dimensionality Reduction	23
Deviation and Anomaly Detection	24
Supervised Learning	21
Regression and Correlation	25
Classification	26
Self-Supervised Learning	27
Generative AI	27
Reinforcement Learning	28
Training and Evaluation in Supervised Learning	30
Ground Truth Dataset	30
k-fold Cross Validation	32
Working with Data Step by Step	33

Summary	38
Chapter 2: Visualization	39
Background and History	39
Basic Visualizations	41
Pie chart	41
Bar chart and Histogram	41
Line chart	42
Scatter plot	42
Broken axis chart	43
More than One Dimensional Data Visualizations	43
Surface plot	44
Contour plot	44
Area plot	45
Radar chart	46
Heatmap	47
Calendar Plot	48
Timeline plot	48
Box plot and Violin plot	49
Hierarchical Visualizations	50
Dendrogram	50
TreeMap	51
Sunburst	52
Graph and Network Visualizations	52
Node-link Diagram	52
Arc Diagram	53
Chord Diagram (Circular Link)	53
Deviation and Change Visualizations	54
Divergent bar plot	54
Dumbbell plot	55
Waterfall	56
Beeswarm	57
3D Data Visualizations	57
Geographical Map Visualizations	59
Choropleth	59
Bubble Map	60
Customized and Other Multidimensional Visualizations	61
Multi-Dimensional Scaling	62
Summary	63
Further Reading and Watching	64

Chapter 3: Statistics & Probability	67
Concepts and Definitions	68
Variable and Value	68
Random or Stochastic Variable	68
Continuous and Discrete Variable	68
Different Data Types	69
Dependent, Independent, and Control Variables	69
Independent versus Dependent Trials	70
First Insight on the Data & Basic Statistical Concepts	70
Mean	71
Arithmetic Mean	71
Geometric Mean	71
Harmonic Mean	72
Median	73
Mode	73
Variance, Standard Deviation, and Covariance	74
Range and Quartile	75
Whisker plot or Box plot	76
Degree of Freedom	78
Probability	78
Basic Probability Concepts	79
Joint Probability	79
Bayes Rule	80
Probability Density/Mass Functions and Cumulative Distribution Functions	81
Statistical Distribution	82
Normal (Gaussian) Distribution	84
Uniform Distribution	85
Beta Distribution	85
Dirichlet Distribution	87
Binomial Distribution	88
Bernoulli Distribution	90
Geometric Distribution	90
Poisson Distribution	92
Weibull Distribution	94
Power Law (Long Tail), Exponential, and ZipfLaw Distributions	95
Zipf law	96
Pareto principle	96
Chi-Square Distribution	97
Boltzmann Distribution	98

Distribution Check with P-P Plot and Q-Q Plot	100
Expected Value and Expectation of a Function	102
Normalization	102
z-score	104
How Much Data is Enough?	105
Central Limit Theorem (CLT)	106
Law of Large Numbers (LLN)	106
Bias in Sampling	106
Confidence Interval	107
Hypothesis and Significance Tests	110
Hypothesis Error	113
A/B Testing and Significance Test	113
Parametric Significance Tests	114
t-test	115
One-sample t-test	115
Independent t-test	116
Paired t-test	116
ANOVA, MANOVA and ANCOVA	117
One-Way ANOVA	117
Repeated Measure (Dependent) ANOVA	118
Factorial ANOVA	118
MANOVA	118
ANCOVA	119
Non-Parametric Significance Tests	120
Chi-Square Test	120
Test of independence	121
Goodness-of-Fit	123
Kolmogorov-Smirnov	124
Kruskal-Wallis Test	124
Mann-Whitney-U Test	126
Significance Test Error Correction	128
Bonferroni correction	128
Tukey correction	129
Effect size	130
Cohen's d Test	131
Odds Ratio	131
Correlation Coefficients	132
Entropy & Information Gain	134
Entropy	134

Measuring Distances between Distributions	137
KL-Divergence	137
Cross-entropy	138
Jensen Shanon Divergence (JS-Divergence)	138
Probability Estimations	140
Maximum Likelihood Estimation (MLE) Approach	141
Expectation Maximization (EM)	142
Summary	143
Further Readings and Watching	145
Part ii: Unsupervised Learning	
Chapter 4: Clustering	149
Similarity and Dissimilarity	149
Similarity Measurement Methods	152
Euclidean Distance	152
Manhattan Distance	153
Mahalanobis Distance	154
Hamming Distance	156
Levenshtein (Edit) Distance	157
Longest Common Subsequences (LCS) Distance	157
Cosine Distance	157
Jaccard distance or Jaccard Index	159
Dynamic Time Warping (DTW)	160
Graph Similarities	162
Clustering Algorithms	164
Partitioning (k-representative) Methods	164
k-mean	166
k-median	166
k-medoid	156
Density Based Methods	168
Density-Based Spatial Clustering Applications with Noise (DBSCAN)	169
OPTICS	171
Hierarchical Methods	175
Single-Linkage (SLINK)	177
DIANA	179
Large Scale Hierarchical Clustering	182
BIRCH	183
CURE	187
Probabilistic and Fuzzy Clustering	189
Gaussian Mixture Model (GMM)	191

Fuzzy C-Mean Clustering (FCM)	194
Other Types of Clustering	194
Clustering Result Evaluation	197
Intrinsic Cluster Evaluation Methods	198
Elbow Method	198
Within Cluster Sum of Squares Errors (WSS)	199
Silhouette Index	199
Dunn Index	200
Davies-Bouldin Index	200
Extrinsic Cluster Evaluation Methods	201
Purity (Entropy)	201
Rand Index	202
Summary	203
Further Reading and Watching	205
Chapter 5: Frequent Itemset, Sequence Mining and Information Retrieval	207
Basic Concepts	208
Information Retrieval Methods	211
Concepts	212
Hash Structures	213
Hash Table	213
Minwise Independent Permutations Hashing (MinHash)	214
Tree Data Structure	217
Binary Tree	218
2-3 Search Tree	219
B-Tree and B+Tree	220
Red-black Tree	220
Trie and Radix Tree	221
Tree Search Methods	222
BFS and DFS	222
Beam Search	222
Monte Carlo Tree Search	224
Bloom Filters	226
Sliding Windows	228
Sliding Window Types	228
SkipList	230
Frequent Pattern Mining Algorithms	232
Apriori	232
FP-Growth	234
ECLAT	239

Sequence Mining Algorithms	241
Generalized Sequential Pattern (GSP)	242
SPADE	244
FreeSpan	247
PrefixSpan	250
Sequential Data Prediction with Hidden Markov Model	254
Probabilistic Graphical Model	254
Hidden Markov Model Concept	255
Markov Model for Prediction	257
Hidden Markov Model for Prediction	259
Likelihood (Forward Algorithm)	261
Decoding (Viterbi Algorithm)	263
Training (Baum-Welch) Algorithm	264
Summary	266
Further Reading and Watching	267
Chapter 6: Feature Engineering	269
Basic Concepts of Feature Engineering	270
Feature Selection	270
Filtering methods	271
Wrapper methods	272
Sequential Feature Selection and Sequential Backward Selection	272
Genetic Algorithm	274
Embedded methods	276
Feature Generation	276
Feature Engineering for Numerical Data	277
Scaling Data	278
Min-Max scaling	279
Standardization and z-score Standardization	279
L1 and L2 norms	280
The Magic Power of Transformation	280
Logarithmic Transformation	282
Box-Cox Transformation	282
Feature Engineering for Categorical Data	284
One-hot Encoding	284
Dummy Coding	285
Effect Coding	285
Feature Hashing	285
Bin Counting	286
Target Encoding	287

Feature Engineering for Textual Data	288
Bag-of-words	291
Subword Tokenization	291
N-Grams	291
Part of Speech Tagging	292
Word Embeddings and Language Models	292
Training a Word Embedding	294
Word Embedding Models	294
Word2Vec	295
GloVe	296
FastText	297
Theme and Keyword Extraction	297
TF/IDF	297
TextRank	297
Rapid Automatic Keyword Extraction (RAKE)	298
Yet Another Keyword Extractor (YAKE)	298
LLM Keyword Extraction	299
Feature Engineering for Image and Video Data	300
Image Processing Concepts and Components	300
Features Extraction from Image Data	303
Image Feature Extraction Algorithms	305
Harris-Stephens (HS) corner detection	305
Maximally Stable Extremal Regions (MSER)	305
Histogram Oriented Gradient (HOG)	307
Scale Invariant Feature Transform (SIFT)	308
Watershed Transformation	312
Video Feature Extraction	314
Motion Vectors	314
Optical Flow	314
3D Convolutional Neural Network (3D CNN)	315
Graph Convolutional Networks	316
Feature Engineering for Signals and Time Series	316
Time series Features	317
Stationarity and non-stationarity time series	317
Periodicity and Seasonal Behaviors	317
Trends	318
Motif	318
Lag	319
Step, Burst and Change Point	319

Signal Features	320
Signal Concepts	320
Types of Signals	320
Signal and Time Series Smoothing	322
Summary	323
Further Reading or Watching	326
Chapter 7: Dimensionality Reduction and Data Decomposition	327
Dimensionality Reduction Methods	328
The Curse of Dimensionality	329
Linear Dimensionality Reduction Methods	330
Principal Component Analysis (PCA)	330
Incremental PCA	334
Linear Discriminant Analysis	335
Fisher Linear Discriminant (FLD)	335
Non-linear Dimensionality Reduction Methods	338
Locally Linear Embedding (LLE)	339
t-distributed Stochastic Neighbor Embedding (tSNE)	341
Uniform Manifold Approximation and Projection (UMAP)	344
Signal and Time Series Decomposition	347
Signal and Time Series Concepts	347
Fourier Transform	349
Wavelet Transform	350
Continuous and Discrete Wavelet Transform	352
Approximate Aggregation Methods (Time Series)	354
Matrix Decomposition	355
Cholesky Decomposition	356
Non-negative Matrix Factorization (NMF)	357
Singular Value Decomposition (SVD)	357
Topic Modeling (Clustering with Matrix Decompositions)	358
Latent Semantic Indexing (LSI)	359
Latent Dirichlet Allocation (LDA)	361
Tensor Decompositions	363
Tensor Concepts	364
Tensor and Matrix Operations	366
Mode-n unfolding	367
n-mode product (multiplication)	368
Tensor Decomposition Methods	369
Canonical or Parallel Factor Decomposition (CP)	369
Tucker Decomposition (TD)	370

Tensor Train Decomposition (TT)	371
Summary	372
Further Reading and Watching	374
Part iii: Supervised Learning	
Chapter 8: Regression, Regularization, and Optimization	377
Objective, Cost, and Loss	378
Linear Regressions	379
Univariate Linear Regression	379
Model Parameters (Coefficients) Estimation	381
Multiple Linear Regression	383
Deciding About Model Variables?	384
Linear Regression Challenges and Resolutions	385
Polynomial Regression	386
Model Parameters (Degrees and Coefficients) Estimation	387
Coefficients	387
Degrees	388
Piecewise, Segmented, or Non-Additive Regression	388
Evaluating the fitness of training sets in linear models.	390
Residual Standard Error (RSE)	391
R^2	391
Root Mean Square Error (RMSE)	392
AutoRegressive Integrated Moving Average (ARIMA)	393
Extrapolation and Interpolation	394
ARIMA Model	394
ARIMA Parameters Estimation	395
Determining the best order of differencing (d)	398
Determining the best order of Auto Regression (p)	398
Determining the best order of Moving Average (q)	398
Logistic Regression	399
Model Parameters Estimation	401
Softmax Regression (Classifier)	403
Model Parameter Estimation	406
Evaluating Regression Models Fitness	407
k-fold cross validation	408
Learning Curve	408
ROC Curve	409
Pseudo R^2	410
Wald Test	410
Information Criterion	411

Likelihood Ratio Test	412
Devils of Model Building	414
Overfitting and Underfitting	414
Bias-Variance Tradeoff	415
Regularization	418
Ridge	419
LASSO	421
Elastic Net	422
Non-Negative Garrote	422
Optimization Algorithms	424
Mathematical Concepts Required for Optimization	424
Derivative	425
Second-Order Derivative	428
Partial Derivative	428
Gradient	429
Jacobian	430
Hessian	430
Integral	430
Taylor Series	432
What is Optimization in Machine Learning?	432
Gradient Descent	434
Types of Gradient Descent	437
Batch Gradient Descent (BGD)	437
Stochastic Gradient Descent (SGD)	438
Mini Batch Gradient Descent (miniBGD)	439
Newton Method	440
Early Stopping	443
Summary	444
Further Readings and Watching	446
Chapter 9: Classification	449
Rule Based Classifier	450
Naive Bayes	451
Bayes Theorem	451
Prediction with Naive Bayes	453
Gaussian Naive Bayes	455
Naive Bayes Prediction Example	456
k Nearest Neighbor (kNN)	459
Voronoi Tessellation	460
KD-Tree	461

Locality Sensitive Hashing (LSH)	463
Support Vector Machine (SVM)	466
Handling Non-linear Data with Maximum Margin Classifier	470
Kernel Trick	471
Kernel Functions	472
Multi-label classification for SVM	473
How Does the SVM Perform the Prediction?	474
SVM Computational Complexity	474
Decision Trees	475
Iterative Dichotomiser 3 (ID3) 6	477
Chi-square Automatic Interaction Detector (CHAID)	479
C4.5	481
Classification and Regression Trees (CART)	484
Decision Tree Challenges and the Need for Pruning.	485
Decision Trees Computational Complexity	487
Ensemble Learning Methods	489
Bootstrap Aggregating (Bagging)	489
Boosting	491
Stacking	492
Random Forest	493
Adaptive Boosting (AdaBoost)	496
Gradient Boosting Decision Tree	501
Regression with GBDT	502
Classification with GBDT	505
eXtreme Gradient Boosting (XGBoost)	510
LightGBM	518
CatBoost	519
How to Select the Best Classification Model?	522
Summary	523
Further Reading or Watching	524
Part iv: Deep Learning	
Chapter 10: Neural Networks and Deep Learning	531
Universal Approximation Theory	533
Biological Neural Network	533
Artificial Neural Network	535
Perceptron Algorithm	539
Multilayer Perceptron	541
Activation Functions	544
Neural Network Cost Functions	547

Quadratic Cost (Mean Squared Error)	547
Cross Entropy	548
Kullback-Leibler Divergence	548
Hellinger (Bhattacharyya) Distance	549
Neural Network Optimizers	551
Stochastic Gradient Descent with Momentum	552
Exponentially Weighted Average	553
SGD with Momentum	554
Nesterov Momentum	555
Adagrad	557
RMSprop	559
Adam	560
Optimizers Summary	561
Backpropagation	562
Forward and Backward Pass Example	567
Regularization in Neural Network	568
Vanishing and Exploding Gradients	568
Weight Initialization	569
Batch Normalization	570
Gradient Clipping	571
Other Normalization Techniques	571
Dropout	572
Early Stopping	573
Convolutional Neural Network (CNN)	574
Convolution and Cross correlation	576
CNN Architecture	577
Different Types of Convolutions	584
Recurrent Neural Network (RNN)	587
Long-Term Short-Term Memory (LSTM)	590
Gated Recurrent Unit (GRU)	593
Bidirectional RNN	594
Deep RNN	596
RNN Examples	596
Summary	599
Further Reading or Watching	601
Chapter 11: Self-Supervised Neural Networks	603
Representation Learning Concepts	605
Generative vs. Discriminative Model	605
Deterministic vs. Stochastic Model	606

Self Organizing Maps (SOM)	607
Boltzmann Machines	611
Restricted Boltzmann Machine (RBM)	612
The Intractable Problem of Z in RBM	613
RBM Training	615
Deep Belief Network and Deep Boltzmann Machine	617
Autoencoders	621
How do Autoencoders work?	622
Autoencoders can Cheat	623
Autoencoder Types	623
Sparse Autoencoder	624
Denoising Autoencoder	626
Contractive Autoencoder	627
Stacked Autoencoder	628
Variational Auto-Encoder (VAE)	630
The Magic of Generative Models in Data Reconstruction	630
How does VAE work?	631
VAE Cost Function	633
U-Net	633
Generative Adversarial Network	637
Training GAN	639
GAN Cost Function	640
GAN Challenges	641
Loss Oscillation	641
Slow Convergence	641
Mode Collapse	641
Uninformative Loss	642
Evaluating GAN Result	642
Inception score (IS)	642
Fréchet Inception Distance (FID)	643
GAN Architectures	644
Conditional GAN (CGAN)	644
Deep Convolutional GAN (DCGAN)	645
Wasserstein GAN (WGAN)	646
Wasserstein or Earth Mover Distance	647
Continuity Condition	649
WGAN Cost Function	649
Weight Clipping	650
WGAN with Gradient Penalty (WGAN-GP)	650

Pix2Pix	651
CycleGAN	653
CycleGAN Architecture	653
CycleGAN Cost Functions	655
Generator and Discriminator Networks	655
StyleGAN Models	656
Noise Sources and AdaIN	657
Progressive Growing	658
Style Mixing	659
StyleGAN Successors	660
Contrastive Representation Learning	663
Contrastive Learning Loss Functions	664
Contrastive Loss	665
Triplet Loss	665
Siamese Network	666
Training Siamese Network	666
Data Preparation for Training	667
Testing Siamese Network	667
Text-to-Image Models	668
Text-to-Image Related Concepts	668
Zero-shot Learning	668
Autoregressive Models	669
Diffusion Models	669
Inpainting and Outpainting	671
CLIP (Contrastive Language–Image Pre-training)	671
Contrastive Pretraining	672
Training	673
Experimenting Zero-Shot learning	673
VQ-GAN (Vector Quantized GAN)	674
VQGAN-CLIP	677
DALL-E Models	677
DALL-E v1	678
DALL-E v2	678
DALL-E v3	680
Imagen	681
Parti	682
Stable Diffusion Models	684
Latent Diffusion Model	684
Training LDM	685

Stable Diffusion XL	686
Stable Diffusion 3	687
Other models that we didn't explain	689
Summary	690
Further Reading and Watching	691
Chapter 12: Deep Learning Models and Applications (Text, Vision, and Audio)	693
Attention and Transformer Neural Network	694
Sequence-to-Sequence	694
Seq2Seq with Attention	696
Attention Models	699
Dot Product Attention	699
Bahdanau (Additive) Attention	699
Luong Attention	700
Self-Attention	700
Cross-Attention	704
Transformer Network	705
Skip Connection and Positional Encoding	706
Skip (Residual) Connection	706
Positional Encoding	706
Multi Head Attention	708
Transformer Architecture	710
Encoder	710
Decoder	710
Masked Multi-Head Attention	712
Structured State Space Sequence Models (S4)	713
State Space Model	713
Linear Time Invariants SSM	714
Structured State Space for Sequence Modeling (S4)	714
Mamba Architecture	714
Large Language Models	717
BERT	718
BERT Architecture	719
BERT Input	719
Pre-training BERT	720
Fine-tuning BERT	721
Derivations of BERT	722
RoBERTa	723
DistilBERT	723
XLM	724

Transfer Learning with a Unified Text to Text Transformer (T5)	724
Generative Pre-trained Transformer (GPT) Models	725
GPT-1	725
GPT-2	726
GPT-3	727
Instruct GPT (Commercialized as ChatGPT)	729
Post ChatGPT LLMs	730
LLama Models	731
LLama2	732
LLama 3	733
Successors of Llama	733
Mistral Model Series	735
Mistral 7B	735
Mixtral-8x7B	736
Mixtral-8x22B	738
Fine-Tuning of LLM	738
Reinforcement Learning from Human Feedback (RLHF)	738
Direct Preference Optimization (DPO)	739
Task-specific and Domain-specific Fine-Tuning	740
Adversarial Fine-Tuning	740
Low-Rank Adaptation (LoRA)	741
Natural Language Model Evaluations	741
LLM Benchmarks	746
Summary of LLMs	747
Computer Vision	749
Image Classification with Convolutional Neural Networks	749
LeNet	750
AlexNet	750
VGG	751
GoogleLeNet / Inception Net	753
Residual Networks (ResNet)	754
Inception-v4 and Inception-ResNet	757
Vision Transformers	758
Object Detection	761
Object Detection Concepts	762
R-CNN Models	763
R-CNN	763
Fast R-CNN	764
Faster R-CNN	765

Single Shot MultiBox Detection (SSD)	767
YOLO Models	771
YOLO v1	771
YOLO v2	772
YOLO 9000	773
YOLO v3	773
Semantic Segmentation and Instance Segmentation	775
U-Net	776
Fully Convolutional Network (FCN)	776
Mask R-CNN	778
DeepLab Models	779
DeepLab v1	779
DeepLab v2	781
DeepLab v3	781
DeepLab v3+	782
Segment Anything Model (SAM)	783
SAM v1	783
SAM v2	787
3D View Synthesis	788
Concepts	789
Neural Radiance Field (NeRF) Model	792
3D Gaussian Splatting	793
Other computer vision models that we didn't explain	797
Audio	800
Audio Signal Concepts and Features	801
Sound Modeling	804
WaveNet Models	805
WaveNet 1	805
Parallel WaveNet (WaveNet 2)	808
Tacotron Models	809
Tacotron 1	809
Tacotron 2	811
wav2vec Models	811
wav2vec v1	811
wav2vec v2	813
Whisper	814
Training Whisper	814
Whisper Architecture	815
Summary	816

Further Reading and Watching	816
Part v: Reinforcement Learning	
Chapter 13: Reinforcement Learning	821
Reinforcement Learning Concepts	822
Exploitation versus Exploration	824
Markov Decision Process	825
Grid-world	826
Deterministic, Nondeterministic, and Stochastic Models	829
Policy	830
Expected Return (Goal)	830
Value, Value function, and Action-value Function	831
State Transition Diagram	832
Bellman Equation	834
Bellman Equation Examples	835
Multi-Armed Bandits Problem	839
Epsilon Greedy (-Greedy)	840
Upper Confidence Bound (UCB)	841
Thompson Sampling	843
Real-world examples of MAB problem	845
Optimal Policy and Optimal Value Function	847
Control and Prediction Problems	848
Tabular Methods	850
Monte Carlo Method (Policy Evaluation and Policy Improvement)	852
Monte Carlo Challenges	852
Monte Carlo Policy Evaluation and Policy Improvement	853
When to use MC method	854
Dynamic Programming (Policy Iteration and Value Iteration)	854
How does a DP algorithm work?	855
Converting Bellman equation into interactive updates	856
When to use DP method	858
Temporal Differences (TD) Learning	860
SARSA and Q-Learning	861
SARSA	862
Q-Learning	863
Dyna-Q	864
Approximation Methods and Deep Reinforcement Learning	867
n-Step Methods and $TD(\lambda)$	867
$TD(\lambda)$	868
Eligibility trace	870

Continuous Space and Function Approximation	872
Deep Q-Networks (DQN)	873
DQN cost function	874
Two problems of Deep RL (i.i.d. samples and non-stationary target)	875
Experience Replay	876
DQN's target network	877
Convolutional Layers of DQN to play Atari games	878
Double Deep Q-Networks (DDQN)	879
Dueling Networks	881
Prioritized Experience Replay (PER)	882
Policy Gradient Methods	884
Policy Gradient Advantages	884
Policy Gradient Objective Function	884
REINFORCE	885
REINFORCE with Baseline	886
Trust Region Policy Optimization (TRPO)	888
Trust Region	888
Minorization-Maximization (MM) algorithm	889
Surrogate Loss	890
Fisher Information	891
Natural Policy Gradient	891
TRPO Algorithm	895
Proximal Policy Optimization (PPO)	895
Penalty Coefficient	896
Clipped Objective	896
PPO as Actor-Critic Model	897
Actor-Critic Methods	898
Asynchronous Advantage Actor-Critic (A3C)	899
Deep Deterministic Policy Gradient (DDPG)	901
Twin Delayed Deep Deterministic Policy Gradient (TD3)	903
Soft Actor-Critic (SAC)	905
Entropy Regularization (Entropy Bonus)	906
Soft Q-value	906
Critics and Target Critics	907
Actor	907
Policy Update and Reparameterization Trick	908
SAC Architecture	909
Dreamer Models	910
Straight-Through Gradient	910

Dreamer v1	911
Dreamer v2	913
Dreamer v3	914
Summary	916
Further Reading or Watching	918
Part vi: Other Concepts and Algorithms	
Chapter 14: Making Lighter Neural Networks and Machine Learning Models	923
Data Compression Algorithms	924
Byte Pair Encoding (BPE)	925
Bitmap index	926
Huffman Coding	927
Lempel–Ziv–Welch (LZW) Encoding	928
Sparse Coding	930
Quantization	933
Vector Quantization and Cluster Quantization	933
Non-uniform Quantization (Signal Companding)	935
μ -Law	936
A-Law	936
Quantization for Neural Networks	937
Fixed Point versus Floating Point Forms	937
How to Quantize?	938
When to Quantize a Neural Network?	939
Dynamic versus Static Quantizations	939
Post-Training Quantization (PTQ)	939
Pre-training Quantization and Quantization-Aware Training (QAT)	940
Channel-wise Quantization	940
Pruning and Sparsification in Neural Network	942
What to prune/sparsify?	943
When to prune?	944
Structured versus Unstructured pruning	944
How to identify pruning candidates?	945
Optimal Brain Damage (OBD) and Optimal Brain Surgeon (OBS)	945
Magnitude-based Pruning	947
Activation-based Pruning	947
Taylor Expansion based Pruning	948
Low-Rank Adaptation of Large Language Models (LoRA)	950
Lighter Self-Attention Mechanisms	951
FlashAttention Models	951
Multi Query Attention and Group Query Attention	953

Multi Query Attention (MQA)	953
Group Query Attention (GQA)	955
Sliding Window and Dilated Sliding Window Attentions	955
Memory Efficient Gradient Descent	956
Gradient Accumulation	956
Gradient Checkpointing	957
Conjugate Gradient	958
Computing the next step in CGD	959
Neural Network Training Paradigms	959
Transfer Learning	960
Multitask Learning	961
Meta-learning	963
Model-based Meta-Learning (MML)	964
Metric-based Meta-Learning	964
Optimization-based Meta-Learning	965
Curriculum Learning	966
Federated Learning	967
Knowledge Distillation	968
Knowledge Distillation Concepts	969
Soft and hard targets	969
Softmax temperature	971
Transfer set	971
Balancing factor	971
Knowledge Distillation Training	971
Common Knowledge Distillation Architectures	972
Automatic Machine Learning	974
Hyperparameter Tuning	974
Random Search	974
Grid Search	975
Neural Architectural Search	975
Summary	976
Further Readings and Watching	976
Chapter 15: Graph Mining Algorithms	979
Basic Concepts	980
Minimum Spanning Tree	981
Prim	982
Kruskal	984
Shortest Path Finding	985
Dijkstra	985

A* Search	987
Matching Algorithm	990
Hungarian Algorithm	990
Hopcroft-Karp Algorithm	992
Centrality Measurement Algorithms	995
Degree Centrality	995
Closeness Centrality	996
Betweenness centrality	997
Eigenvector centrality	998
Link Analysis Algorithms	998
PageRank	999
Hyperlinked Induced Topic Search (HITS)	1002
Community Detection in Graphs	1004
Spectral Clustering	1005
Louvain	1007
Leiden	1010
Graph Neural Network	1011
Challenges of working with Graphs in Neural Network	1012
Message Passing Layer (Propagation)	1013
Graph Pooling	1014
Spectral versus Spatial Representation	1015
Graph Convolutional Network (GCN)	1015
Graph Attention Network	1001
Graph SAGE	1017
High-Dimensional Search with Graphs	1020
Hierarchical Navigable Small World	1021
Summary	1023
Further Reading and Watching	1024
Chapter 16: Challenges of Working with Data	1025
Problem Complexity	1025
Sampling from Complex Data Structure	1028
Stratified Sampling and Clustered Sampling	1028
Monte Carlo (MC) Sampling	1029
Monte Carlo Markov Chain (MCMC) Methods	1031
Metropolis-Hastings	1032
Gibbs Sampling	1033
Importance Sampling	1036
Noise	1037
Types of Noises	1021

Noises Based on Statistical Distributions	1038
White noise	1038
Gaussian noise	1038
Other types of noises	1039
Salt-and-Pepper noise	1039
Gradient noise	1039
Speckle noise	1040
Noise Reduction	1040
Machine Learning based noise reduction	1041
Signal-based Noise Filtering	1041
Butterworth filter	1043
Wiener filter	1043
Kalman filter	1044
Imbalance Data Problem and Data Augmentation	1049
Data Augmentation Methods	1050
Tabular Data	1050
Image	1051
Text	1053
Signal and Time Series	1054
Reconstructing Missing Data	1055
Imputation	1056
Mean/Median or Mode	1056
K-Nearest Neighbors (kNN) for Imputation	1056
Hot deck imputation	1056
Cold deck imputation	1056
Regression Imputation	1057
Griffin-Lim Algorithm (GLA)	1057
Customized methods	1058
Interpolation	1058
Linear and Bilinear interpolation	1058
Polynomial interpolation	1060
Spline interpolation	1060
Kriging interpolation	1060
Radial Basis Function Interpolation (RBF)	1062
Anomaly and Outliers Detection	1062
Unsupervised Methods	1063
Isolation-based Methods	1063
Isolation Forest (iForest)	1063
One-Class SVM	1065

Local Outlier Factor (LOF)	1065
Random Sample Consensus (RANSAC) Algorithm	1066
Data and Model Changes over Time	1067
Concept, Model and Data Drift	1068
Tackling Drift Challenges	1068
Cold Start Problem	1069
Tackling Cold Start Problem	1070
Raters Agreement Methods	1071
Percentage Agreement	1071
Cohen's Kappa	1071
Fleiss' Kappa	1072
Krippendorff's Alpha	1073
Utilizing LLMs	1075
Retrieval Augmented Generation (RAG)	1075
Fine Tuning the LLM with External Data	1075
Indexing the External Data	1075
Store External Data in an Embedding	1076
Prompt Engineering	1077
Zero-shot and Few-shot Prompting	1077
Chain-of-Thought (CoT)	1078
Self-Consistency	1079
Tree-of-Thought (ToT)	1081
Fairness, Bias, and Transparency	1083
Concepts and Terms	1084
Fairness Parity Metrics	1085
Bias and Fairness Mitigation Methods	1086
Interpretability and Transparency	1088
SHapley Additive exPlanations (SHAP)	1088
Local Interpretable Model-agnostic Explanations (LIME)	1091
Summary	1092
Further Reading and Watching	1093
Bibliography	1095
Index	1123